



NTNU – Trondheim
Norwegian University of
Science and Technology

The advection equation using six-function MPI


The example code archive

- I put another copy of the sequential advection code into “02_advection_sequential”
- It’s pretty much the same thing as before, but I have
 - Divided the code into (hopefully) aptly named functions, and
 - Increased the problem size in order to make it run for a little bit
 - If you want to adjust the size yourself, mind that different sizes of output files require a modification to the plotting script, the number of elements to plot is hardcoded in there
- The remaining two directories contain
 - A partial version that can only initialize the program and save its state
 - A complete version that includes the numerical solver loop



We need to split up the work

- We have a long, linear array, and some number of workers to employ
 - Let's draw 4, just as an example
- Here's a popular, but not-so-good solution
 - Allocate full array for everyone, but just work on part of it



- Add up everyone's partial solutions at the end







Why would anyone do that?

- It's really simple to work out array indices when everyone has the same coordinate space
 - You'll see in a minute
- It's still not a great idea, though
 - It limits the maximum problem size to the amount of memory 1 rank can allocate
- I am not going to say that it's bad in every context
 - Small problems are also worth solving
- I am going to say that it's an impediment to scalability
 - So there



Another way to split work

- Divide the problem size by the rank count, and allocate separate parts


Rank 0		Global elements 0-3, local indices 0-3
Rank 1		Global elements 4-7, local indices 0-3
Rank 2		Global elements 8-11, local indices 0-3
Rank 3		Global elements 12-15, local indices 0-3

- We get to concatenate these when saving the state of the entire domain



From the boundary line

- The way we allocated/indexed the array in the serial version, we added 2 extra points for the boundary condition
...and gave them indices -1 and N...
- Let's do that everywhere here, too




- Only $U(-1)$ by rank 0 and $U(4)$ by rank 3 will *actually* represent the problem domain's boundary
- We'll have use for the others as well



Initialization

- Since we've split the coordinates,
 - Rank 1 must know that its element 0 is global element 4
 - Rank 2 must know that its element 0 is global element 8
 - *Etc.*



- One possible solution:

```
int_t my_origin = rank * (N / size);
```
- This requires all parts to be equally large



Initialization


- What if the domain size isn't divisible by the number of ranks?
- There are three schools of thought:
 - Stop the program, and demand a particular problem-size / rank count relationship
 - Give the last rank less work (either in a smaller allocation, or padding out the domain data with zeros at the end)
 - Give 1 extra element to a suitable subset of the ranks



Initialization

This week's example code goes for the last option:

```
local_sizes = malloc ( size * sizeof(int_t) );
for ( int_t r=0; r<size; r++ )
    local_sizes[r] = (int_t)( N / size ) + ((r<(N%size)) ? 1:0);
```



← (Example: size 18 problem with 4 ranks)

- The result is that every rank gets an array with the others' subdomain sizes in it
 - Per the illustration, `local_sizes` would contain [5, 5, 4, 4]



Initialization

- Armed with the knowledge of how big the preceding problem parts are, each rank can calculate what its own origin index corresponds to globally

```
int_t my_origin = 0;  
for ( int_t i=0; i<rank; i++ )  
    my_origin += local_sizes[i];
```

- A small amount of extra typing, but the code only has to run once, and it's very short
- Now that we can calculate x-positions from the indices, we can plug in the function that creates the initial advection state
- Each rank can set up its part separately



PROTIP:

A picture tells a 1000 words

- The first thing I do when starting a parallel program is to invent a way to draw pictures of the global state
 - Parallel programs run in a mish-mash order that can be different every time you launch the program, so debugging with print statements gets messy
 - It's more feasible if you make every rank write in a separate file, but that still makes it hard to see the interplay between them
- When doing this, it's important to make double-triple sure that your visualization actually matches the program state
 - Bugs that create inaccurate pictures come back to haunt you later

Saving global state

- The files we used in the sequential code are just a long list of floating point numbers stored in binary
- To make that again, we'll need to concatenate the numbers from all ranks, in rank order
- We can nominate rank 0 to be our “I/O-master”, who can collect all the parts and put the file together



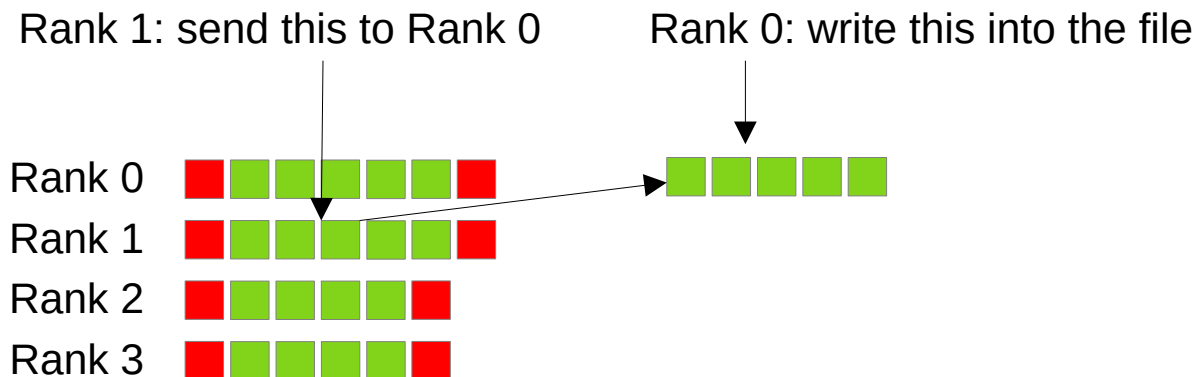
Rank 0 needs an extra buffer

- Because of the way we partitioned, rank 0 will always have (one of) the biggest subdomains, so we can use its size
- Step 1:



Rank 0 needs an extra buffer

- Rank 0 can now loop over the remaining ranks, and receive their sub-domains in the buffer
- Step 2:



- Steps 3 and 4 are just like step 2



This is a bottleneck

- We're serializing the execution, rank 0 has to wait for all the ranks in turn, and do something sequential for them
- Another way would have been to let the ranks take turns to open the file and append to it
 - Still sequential, but with less communication
- We could make each rank save its own file, and concatenate them after the program has finished
 - Parallel, but creates more logistics afterwards
- Yet another would be to have everyone write at the same time
 - But we're only doing 6-function MPI today
- Saving doesn't happen on every iteration anyway



The example code archive

- This is the state of the code in the subdirectory “03_init_and_cleanup”
- It divides the problem and makes allocations
- It initializes all the arrays
- It saves the global initial state in a file
- It releases all the arrays again



Adding the solver

- If we draw the arrays of the ranks side-by-side, we can see an issue with the numerical method:
 - The boundaries are at ranks 0 and 3, but the calculation needs two neighbor values everywhere



Rank 0



Rank 1



Rank 2




Rank 3




The value of good neighbors

- The last value by rank 1 needs the first value from rank 2 as a neighbor
 - We can send it a copy




- The first value by rank 2 needs the last value from rank 1 also
 - We can send copies it in the other direction as well



Border exchange

- This operation is common, and it is called a **border exchange**
- The artificial extra-points are often called **ghost points**, together they are often referred to as a subdomain's **halo**
- Since we're using a periodic boundary, border exchange takes care of that too, as long as we connect the first and last ranks:


```
left_neighbor = ( rank + size - 1 ) % size;
right_neighbor = ( rank + size + 1 ) % size;
```
- Adding the extra 'size' here is just because moduli of negative numbers aren't a thing in C



Adding the solver

- When we've taken care that all the surroundings of the solver are as it expects, it can simply be used the way it was
- Only difference is that its loop has to go from 0 to the rank's subdomain size, instead of to N
- That's the code, let's see if it runs any faster...

