



NTNU – Trondheim
Norwegian University of
Science and Technology

Introduction to OpenMP

Open Multi-Processing

- OpenMP is almost as storied as MPI
 - Version 1.0 was published in 1997
- It occupies a similar position as a *de facto* standard tool for scientific number crunching
- It targets shared-memory systems
 - MPI's unit of parallelism is the process
 - OpenMP's unit of parallelism is the thread
- You can think of it as a more convenient way to handle pthreads
 - It's not *obliged* to be implemented using pthreads, but it very often is
 - It doesn't *only* encompass threads, but that is by far its most common use



What's in a name?

- Many moons ago, a major MPI implementation called LAM/MPI merged with two less prominent ones
- A new name was required after the merger, so they settled on...
“*OpenMPI*”
 - **sigh**
 - I think they should have chosen a different name, but here we are
- OpenMPI is one out of many MPI implementations
- **OpenMP** is an entirely different programming model, with its own specifications document
 - Various implementations of this interface emerge and disappear...
 - So it goes.



Parts of the puzzle

- As we've seen
 - MPI is a separate 3rd party library of functions that you install in addition to your compiler
 - Pthreads are provided by the operating system, the function declarations come with the compiler
- OpenMP is a little bit of both
 - Its core is a set of language extensions that must be supported by the compiler
 - It also has a runtime library of functions that you can call to inspect the state of what the compiler has generated



Language extensions?

- Yes; C has a standardized way to do nonstandard things, so to speak
- The `#pragma` directive can be followed by some text that the compiler will discover during its initial scan of the program code
 - If it understands the text, it can insert appropriate code to replace the directive with
 - If it doesn't understand the text, the compiler is free to discard it
- This way, compilers can support optional features in the code that
 - Work when you use a compiler that supports them
 - Don't break the program even if you use a compiler that doesn't support them



A hypothetical example

- `#pragma` can ask for literally anything:

```
#include <stdio.h>
int main() {
    printf ( "Hello, world!\n" );
    #pragma play me a song
    return 0;
}
```

- You can compile this code without issue (try it at home)
 - My compiler only makes the usual hello-world binary without any special effects
 - It still reads the command
 - It just doesn't know what to do with it, and throws it away
- Given a compiler that supported it, this directive could produce a musical executable



A more practical use

- Pthreads code is tediously repetitive
- We have to do the same things over and over:
 - Declare, initialize, use, and destroy a mutex for every thing that needs protection
 - Declare, initialize, use, and destroy a cond for every signal
 - Declare, initialize, use, and destroy an object for every barrier
 - Simple sets of operations make for lots of repetitive typing
- Since the code is practically the same over and over, we might as well make the compiler generate it
- It can figure out what to generate from a tiny language embedded in well-placed `#pragma` directives
- That's OpenMP's mechanism of choice



Stack contexts

- We've covered how a function call encapsulates a local set of values on the call stack
 - That's the connection between function calls and pthread creation
- Other local scopes also contain stack contexts
- Consider this program fragment:

```
int main ()
{
    int a = 42, b = 32, c = 0;
    {
        int a = 64;
        c = a - b;
    }
    printf ( "a = %d, b = %d, c = %d\n", a, b, c );
    return 0;
}
```

The output is "a = 42, b = 32, c = 32"



What's going on?

- An open `{ /* basic block */ }` establishes a local stack context
 - Just like a function call, except that it doesn't have arguments and return value
- A basic block can appear wherever a statement can
 - That's how we make if-branches and loop bodies
(and function bodies, for that matter)

```
int main ()
{
    int a = 42, b = 32, c = 0;
    {
        int a = 64;
        c = a - b;
    }
    printf ( "a = %d, b = %d, c = %d\n", a, b, c );
    return 0;
}
```

Basic block acting
as a statement



NTNU – Trondheim
Norwegian University of
Science and Technology

Block local scope

- Even when they don't have names and arguments, basic blocks let you declare variables that live only inside the block
- That is a stack context at work:

Execution
is here
→

```
int main ()  
{  
    int a = 42, b = 32, c = 0;  
    {  
        int a = 64;  
        c = a - b;  
    }  
    printf ( "a = %d, b = %d, c = %d\n", a, b, c );  
    return 0;  
}
```

Stack state

c=0
b=32
a=42



NTNU – Trondheim
Norwegian University of
Science and Technology

Block local scope

- After a few more steps, another stack context has been started
- We now have two variables called 'a'
 - The most recent one is near the top of the stack in the scope of the most recently opened block
 - The other one sits in the stack space of the enclosing block

```
int main ()
{
    int a = 42, b = 32, c = 0;
    {
        int a = 64;
        c = a - b;
    }
    printf ( "a = %d, b = %d, c = %d\n", a, b, c );
    return 0;
}
```

Execution
is here



Stack state

a=64
c=0
b=32
a=42




NTNU – Trondheim
Norwegian University of
Science and Technology

Block local scope

- When the time comes to evaluate this expression
 - The nearest declaration of a is used
 - The current block's context doesn't contain b and c, so they are tracked down in the enclosing scope
 - (If they hadn't been there, the next thing would be to check if they were declared globally)

```
int main ()
{
    int a = 42, b = 32, c = 0;
    {
        int a = 64;
        c = a - b;
    }
    printf ( "a = %d, b = %d, c = %d\n", a, b, c );
    return 0;
}
```

Execution
is here



Stack state

a=64

c=32

b=32

a=42



NTNU – Trondheim
Norwegian University of
Science and Technology

Block local scope

- When the block ends, its local context is deleted from the stack
 - the “old” value of a becomes the topmost one in our stack context again
 - Hence, c is 32 even though a-b is 10 now
 - We temporarily created a stack context with a different local variable in

```
int main ()
{
    int a = 42, b = 32, c = 0;
    {
        int a = 64;
        c = a - b;
    }
    printf ( "a = %d, b = %d, c = %d\n", a, b, c );
    return 0;
}
```

Execution
is here
→

Stack state

c=32

b=32

a=42



NTNU – Trondheim
Norwegian University of
Science and Technology

Stack contexts can be threads

- We might as well leave it to the compiler to write the thread spawning and joining logic
- There's a program called 'hello_openmp' in today's example archive
- Notice that the Makefile has added the flag

-fopenmp

to the C compiler's command line

- This enables OpenMP using gcc and clang
- Using icc, the flag is -qopenmp
- Using MSVC I don't know what it is, but it's something (read the manual)



We have the magic ingredients again

- Armed with a thread count and a thread id#, we can solve all the embarrassingly parallel problems again
 - Pick a task based on the id#
 - Handle it
- OpenMP has a far richer set of concepts and tools than this
 - So far, it's definitely the least amount of typing to make a hello world example parallel, though

How many threads do we get?

- By default, OpenMP assumes that you want one thread per core that your O/S recognizes
- You can adjust it without recompiling the program
 - if you set the environment variable
`OMP_NUM_THREADS`
in your shell, OpenMP will look it up there
- You can also hard-code it into the program

```
#pragma omp parallel num_threads(4)
```

will always spawn 4 threads, overriding both your system information and the environment variable
 - There's rarely a good reason to do this, though



We can do locking

(just like pthreads)

- The example program 'pi_mutex_omp.c' is (functionally) identical to last lecture's 'pi_mutex_fast' example
 - Computes local estimates per thread
 - Uses a mutex data structure to avoid race conditions for a global value
- There are smoother ways to do this in OpenMP
 - Don't take it as a wonderful implementation strategy
 - I just wanted to demonstrate that OpenMP code can act precisely like pthread code

We can do barriers

(just like pthreads)

- The example program 'pi_barrier_openmp.c' is (functionally) identical to last lecture's 'pi_barrier' example
 - Repeats computation 10 times
 - Synchronizes between repetitions, to avoid race conditions when resetting the global value
- There are smoother ways to do this in OpenMP as well
 - Don't take it as a wonderful implementation strategy
 - I just wanted to demonstrate... oh, you get the point



We can't do pthread_cond_t

- Inter-thread signals aren't a thing in OpenMP
- OpenMP threads aren't supposed to be sleeping, they're supposed to be computing something
 - The constructs contain lots of busy-waiting, critical sections are expected to be as short as possible
 - Oversubscribe thread counts at your own peril
- If you want to yield CPU cores, just shut down the threads instead
 - They're very easy to bring back again

(There is actually a different technique as well, but we'll get back to it later)



How safe is this stuff?

- It is a little easier to write correct OpenMP code than it is to write correct pthreads code
 - That's mostly because it requires you to consider fewer details at a time, though
- The “*gentleman's agreement*” philosophy still applies
 - OpenMP makes threads when you tell it to
 - If you treat a shared variable as if it were private, OpenMP will take you at your word
 - If you say that something should be parallelized when it should not, you will get programs that compute wrong answers



Our reason to do this

- Today's examples are really written in a pretty clunky style
 - It is actually quite rare to need the thread id# and count for anything in OpenMP
 - I just wanted to show you that they are there, so as to demonstrate that the correspondence to pthreads lurks just below the surface
- That's kind of why we covered pthread programming in the first place
 - Like assembly code, it's not very common to need explicit pthread code
 - Like assembly code, it's good to know what's going on even if you don't type it out by hand



Going forward

- Next time, we'll start on the rich library of OpenMP abstractions

