



**NTNU – Trondheim**  
Norwegian University of  
Science and Technology

## **OpenMP: Tasks**

# Worksharing-style OpenMP

- The way we have been partitioning computational workloads so far depends strongly on the layout
  - We expect to have some large, multidimensional arrays
  - Iterations can proceed along each dimension in turn
  - The size of the workload is fixed and known when the loops begin
- This is pretty close to optimal for parallel for-loops
  - Great for linear algebra
  - Great for Fourier transforms

...but not all programs simulate physical phenomena



# Issues with the worksharing view

- Suppose we start a bunch of threads and inside the parallel region, we have a loop that calls a function with an array of points:

```
#pragma omp parallel for
for ( int i=0; i<N; i+=BLOCK )
    do_something ( &(amp;data_points[i]), BLOCK );
```

- How do we deal with parallelism at the receiving end?

```
void do_something ( mytype_t *data_points, int count )
{
    #pragma omp for
    for ( int x=0; x<count; x++ )
        data_points[x] = sqrt(x);           // ...or whatever
}
```

Here we have assumed  
that the threads are  
already online



# Issues with the worksharing view

- Suppose we **don't** assume that the threads have already gone live

```
void do_something ( mytype_t *data_points, int count )  
{  
    #pragma omp parallel for  
    for ( int x=0; x<count; x++ )  
        data_points[x] = sqrt(x);    // ...or whatever  
}
```

- Now we can only call the function from outside of parallel regions

...and get the stop/start-effect of spawning and joining threads for every block we transfer



# Nested parallelism

- Nested parallelism is when one work-package spawns more work-packages that should be distributed among threads
- The worksharing directives aren't very good at this kind of thing
  - They practically assume that all threads will participate in one big loop-schedule that is ready when the loop begins
  - As we know, this is not **100%** necessary  
(we can have multiple loops with no-wait clauses in a parallel region)  
but it's definitely the default assumption
  - Demonstrably clumsy when we want to separate the list of things to do from the team of threads that do them




# Task-based programming

- This is an alternative view of how to express parallelism
- The idea is similar to the original thought behind pthreads:
  - Take a block of work and dispatch it for background execution
  - Record which blocks of work depend on the other ones
  - Assign them to the team of threads in an order that matches their dependencies



# From a bird's eye view

- Task-based programs generate dependency graphs behind the scenes:



- This arbitrary example-graph would suggest that
  - Tasks T1, T2, T3 can run in parallel
  - Outputs from T1,T2 are needed for T4, and outputs from T2,T3 are needed for T5
  - T4, T5 can run in parallel
  - Etc. etc.*


T2,T3 are



NTNU – Trondheim  
Norwegian University of  
Science and Technology

# Worksharing directives can be task graphs too

- Their shape is just very trivial and uninteresting:



*i.e.* loop iterations are independent, and synchronize when they are finished



# OpenMP tasks

- OpenMP admits the creation of arbitrary dependency graphs through the *task* directive

- We can write

```
#pragma omp task
{
    do_some_stuff();
}
```

and the block's context will be whisked away into an internal queue somewhere, to be executed at the first opportunity

- If we want to wait for all spawned tasks to be finished, there is

```
#pragma omp taskwait
```



# Function calls can be tasks

- This maneuver

```
#pragma omp task
```

```
some_useful_function ( arg1, arg2, arg3 );
```

will take the whole function call to 'some\_useful\_function' and make a background task out of it

- We can also declare functions to be tasks by definition

```
#pragma omp task
```

```
void some_useful_function ( int arg1, int arg2, int arg3) { ... }
```

which will task-ify every call we make to it



# Tasks with and without threads

- You can make tasks out of things without having started a parallel region
  - They'll just be added to a list and run in sequence
- When there *is* a live team of threads active, they'll pick up tasks from the task graph in parallel



# The wonderful part of this

- The body of a task is at liberty to create more tasks
  - Their dependencies can be inferred from their arguments and uses of their return values
  - Alternatively, *taskwait* directives, if you want to be explicit about it
- It's not necessary to assume any particular relationship between the thread count and the number of tasks
  - Tasks-spawning-tasks-spawning-tasks can nest as deeply as you like, they will all be run in due time



# Impractical application

- Making tasks out of loop iterations serves little purpose
  - As we have demonstrated, you can do it just fine
  - There's even a directive **#pragma omp taskloop** that automates making a task out of each iteration in a loop
- It doesn't work very well with the loops we've been using worksharing directives on:
  - This only exposes the same amount of parallelism as we did before
  - It comes with the additional overhead of constructing the trivial taskgraph internally



# Practical application

- Tasks come into their own when you're solving problems that are impractical to express as loops
- Divide-and-conquer algorithms are a splendid example
  - i.e. problems where the parallelizable work comes out of each nesting level in a recursive function call:
    - Make a task out of the first call...
    - Spawn a couple of smaller tasks at the 2<sup>nd</sup> nesting level...
    - Make even smaller tasks at the 3<sup>rd</sup> level...
    - Split them up into more tasks at the 4<sup>th</sup> level...
    - ...you get the picture...



# Starting the chain reaction

- With a recursive divide-and-conquer problem, say

```
void here_we_go() {  
    #pragma omp task  
    do_the_first_half();  
    #pragma omp task  
    do_the_second_half();  
}
```

it's natural to try and write

```
#pragma omp parallel  
{  
    #pragma omp task  
    here_we_go();  
}
```

- This is a mistake
  - N threads will start N individual task-trees that all do the same thing



# The common pattern

- If you have a recursive tree of function calls that spawn tasks, the top level tends to look like this

```
#pragma omp parallel  
{  
    #pragma omp single  
    start_the_circus();  
}
```

Prepare some threads to pick up the tasks

One thread starts the recursion, it will soon create enough parallel work for everyone



NTNU – Trondheim  
Norwegian University of  
Science and Technology



# Example time

- We need a divide/conquer type of algorithm
  - This is different from our classical HPC number-crunching applications
- I've gone with *quicksort*
  - You have (supposedly) already encountered this method in Algorithms & Data Structures, but we can repeat it briefly

# A quick review of quicksort

- Pick a range in an unsorted array of numbers

1	6	3	4	14	2	7	8	9	10	11	12	13	5	15	16
---	---	---	---	----	---	---	---	---	----	----	----	----	---	----	----

(This one is only very mildly out of order, in the interest of brevity)

- Choose a *pivot* number
  - Say, 8
- Search from the low end until you find a number  $>$  pivot

1	6	3	4	14	2	7	8	9	10	11	12	13	5	15	16
---	---	---	---	----	---	---	---	---	----	----	----	----	---	----	----

\_\_\_\_\_ ↑ A-ha!

- Search from the high end until you find a number  $<$  pivot

1	6	3	4	14	2	7	8	9	10	11	12	13	5	15	16
---	---	---	---	----	---	---	---	---	----	----	----	----	---	----	----


\_\_\_\_\_ ↑ A-ha!

heim


Norwegian University of  
Science and Technology

# A quick review of quicksort

- When you have found two suitable numbers, swap them:



- When your search-pointers pass each other, the array is a little bit more sorted than it was:



- One part is an unordered list of  $<$ pivot numbers
- The other is an unordered list of  $>$ pivot numbers



# Divide and conquer

- These two parts can now be passed along for further quick-sorting:

1	6	3	4	5	2	7	8	9	10	11	12	13	14	15	16
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----

- Each part will have its own beginning and end  
...and we'll pick a new pivot in the range of each of them
- When this process has reached a list length of 1, a single number is sorted by default



# Choosing pivots

- The algorithm behaves a little bit differently depending on how you choose the pivot
  - I've gone with the MOT (median-of-three) approach:  
Compare the first, last, and middle elements, and use the median
- There are a few different ways to manipulate memory as well
  - The example implementation sorts in-place, i.e. it overwrites the unsorted array with its sorted equivalent



# Choosing programming languages

- Today's example is written in C++
  - Sorry about that
  - The reason is that I wrote it in order to compare OpenMP with another task-friendly programming model that only exists for C++
- Hopefully, you can read it anyway
  - It's not doing anything super object-oriented, functional, or any meta-programming, so it's pretty C-like after all
  - If you feel it is unfair of me to swap languages mid-semester, just tell me, and I will happily translate it into plain C
  - It's not a lot of work, I'm just not eager to rewrite things unless I know they will be useful to someone



# A small disclaimer

- OpenMP implementations differ from compiler to compiler
  - I get speedup out of this implementation with GCC/g++
  - I **don't** get speedup out of it with LLVM/clang++
  - This isn't necessarily universally true, I'm just mentioning it
  - If you're on MacOS, the 'gcc' package may have installed a version of clang masquerading under the name gcc
  - Call it with '--version' if you're uncertain, and see what it answers
  - Don't ask me why the responsible package-manager people have chosen to do this, because I don't understand it

*(If you understand it, I'd love to hear what the reason is)*