

# Chapter 1

## Fundamentals of Quantitative Design and Analysis

*Acknowledgement: Slides are adapted from Morgan Kaufmann companion material*

# The Computer Revolution

- Progress in computer technology
  - Underpinned by Moore's Law
- Makes novel applications feasible
  - Computers in automobiles
  - Cell phones
  - Human genome project
  - World Wide Web
  - Search Engines
- Computers are pervasive


# Classes of Computers

- Internet of Things/Embedded computers
  - Hidden as components of systems
  - Stringent power/performance/cost constraints
- Personal Mobile Devices (PMDs)
  - Wireless devices with multimedia user interfaces such as cell phones, tablet computers, etc.
  - Stringent power/performance/cost constraints
- Desktop computers
  - General purpose, variety of software
  - Subject to cost/performance tradeoff
- Server computers
  - Network based
  - High capacity, throughput, reliability
- Clusters/Warehouse-Scale Computers
  - Collections of desktop computers or servers connected by a network to form a larger computer
  - Price-performance oriented

# Improving Performance Through Exploiting Parallelism

- Exploiting parallelism is the main way of improving computer performance
- Types of parallelism
  - Data-level parallelism (DLP): Same operation, different data
  - Task-level parallelism (TLP): Divide work into tasks that can operate mostly in parallel
- Sources of parallelism in computers:
  - **Instruction-Level Parallelism (ILP):** Independent instructions can be executed in parallel
  - **Vector processing or Single Instruction Multiple Data (SIMD) parallelism:** A single instruction can be applied to multiple data elements at the same time.
  - **Thread-Level Parallelism (TLP):** Software can be organized as different threads that can operate in parallel (commonly threads communicate)
  - **Request-Level Parallelism:** In request-response server systems, software can process tasks in parallel (threads don't explicitly communicate but may access shared data structures)
  - **Memory-Level Parallelism (MLP):** Processor memory requests (loads and stores) can be issued in parallel to hide memory latencies

# Flynn's Taxonomy



- Single instruction stream, single data stream (SISD)
- Single instruction stream, multiple data streams (SIMD)
  - Vector architectures
  - Multimedia extensions to instruction sets (e.g., SSE)
  - Graphics processor units (GPUs)
- Multiple instruction streams, single data stream (MISD)
  - No commercial implementation
- Multiple instruction streams, multiple data streams (MIMD)
  - Tightly-coupled MIMD: multi-cores, many-cores
  - Loosely-coupled MIMD: clusters, data centers

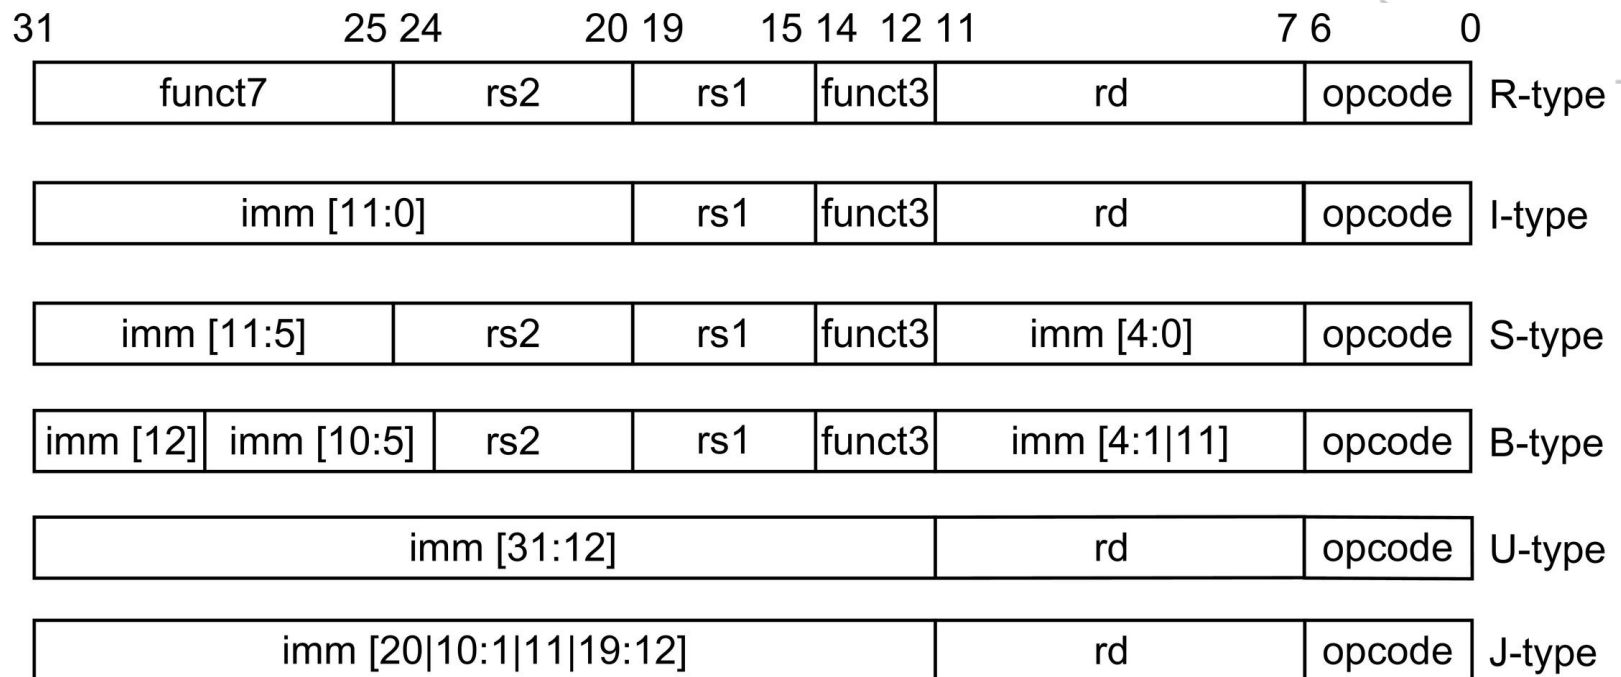
# WHAT IS COMPUTER ARCHITECTURE?

# Classical Computer Architecture

## Definition: The ISA

- The Instruction Set Architecture (ISA)
  - The actual, programmer visible instructions set
  - In other words: The language you can use to get the computer to do stuff
  - We will focus on RISC-V, alternatives are x86, ARM, etc.
- ISA components and choices (more on this later):
  - Operand types (typically registers or memory locations)
  - Memory addressing (typically byte-addressing)
  - Addressing modes (e.g., register, immediate, displacement)
  - Types and sizes of operands (word, char, double word, etc.)
  - Operations (data transfer, logical, control or floating point)
  - Control flow instructions (conditional branches, unconditional jumps, etc.)
  - Encoding (fixed vs. variable length)

# RISC-V Instruction Formats






# Genuine Computer Architecture Definition

- Computer architecture objective: Design the *organization* and the *hardware* to meet objectives and functional requirements
- Computer organization:
  - High-level aspects of a computer's design
  - Examples: Memory system, interconnect, design of the CPU, etc.
- Computer hardware:
  - Detailed logic design and packaging of the computer
- Final definition: Computer architecture covers the ISA, the organization and the hardware.

# Below Your Program



- Application software
  - Written in High-Level Language (HLL)
- System software
  - Compiler: translates HLL code to machine code
  - Operating System: service code
    - Handling input/output
    - Managing memory and storage
    - Scheduling tasks & sharing resources
- Hardware
  - Processor, memory, I/O controllers

# Levels of Program Code

- High-level language
  - Level of abstraction closer to problem domain
  - Provides for productivity and portability
- Assembly language
  - Textual representation of instructions
- Hardware representation
  - Binary digits (bits)
  - Encoded instructions and data

High-level  
language  
program  
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```

Compiler

Assembly  
language  
program  
(for MIPS)

```
swap:
    muli $2, $5, 4
    add  $2, $4, $2
    lw   $15, 0($2)
    lw   $16, 4($2)
    sw   $16, 0($2)
    sw   $15, 4($2)
    jr   $31
```

Assembler

Binary machine  
language  
program  
(for MIPS)


```
000000001010000100000000000011000
000000000000110000001100000100001
100011000110001000000000000000000
100011001111001000000000000000100
101011001111001000000000000000000
101011000110001000000000000000100
00000011111000000000000000001000
```

# Inside the Processor (CPU)

- A processor capable of general computation needs to be able to do arithmetic, conditional branches and memory access
- Arithmetic
  - Datapath: performs operations on data (i.e. arithmetic)
- Conditional branch
  - Control: selects next instruction depending on branch outcome
  - Control: enables the parts of the datapath needed to execute a given instruction
- Memory access
  - Off-chip access is often slow
  - Solution: Cache (Small fast hardware-managed SRAM memory)


# Multi-core Example

- Intel Skylake Core i7 multi-core: 28 processor cores




# TECHNOLOGY TRENDS

# Single-Core Processor Performance Growth






# Latency vs. Bandwidth





# Clock Frequency



# PRINCIPLES OF QUANTITATIVE DESIGN

# Principles of Quantitative Design

- Take advantage of parallelism
  - Exploiting parallelism is the pervasive strategy for improving the performance of computers
- Leverage locality
  - Programs tend to reuse data and instructions – think of loops and arrays
  - This results in reuse over time (temporal locality) and in (memory) space (spatial locality)
- Focus on the common case
  - Computer architecture is full of trade-offs – favoring the frequent case over the infrequent case typically results in better performance
  - Acknowledge Amdahl's law: A massive improvement on an infrequent case gives negligible overall improvement

# Pitfall: Amdahl's Law

- Improving one aspect of a computer and expecting a proportional improvement in overall performance

$$T_{\text{Improved}} = \frac{T_{\text{Affected}}}{\text{ImprovementFactor}} + T_{\text{Unaffected}}$$

- Example: Multiply accounts for 80s out of 100s total runtime
- How much improvement in multiply performance is necessary to get 5X overall speed-up (i.e., reduce runtime to 20s)?

$$\lim_{n \rightarrow \infty} f(n) = \frac{80}{n} + 20 = 20$$





# MEASURING PERFORMANCE

# Understanding Performance

- Algorithm
  - Determines number of operations executed
- Programming language, compiler, and architecture
  - Determine number of machine instructions executed per operation
- Processor and memory system
  - Determine how fast instructions are executed
- I/O system (may include the Operating System (OS))
  - Determines how fast I/O operations are executed

# Defining Performance

- Which airplane has the best performance?



# Response Time

- Time from issuing a command to its completion
  - This is often referred to as the *turn-around time*
- Another response time definition: Time from issue to first response
- *Execution time* is the time the processor is busy execution the program
  - Turn-around time includes the time the process waits to be executed, execution time does not
  - Also: user execution time vs. system execution time



# Response Time and Throughput

- Throughput
  - Total work done per unit time
- How are response time and throughput affected by
  - Replacing the processor with a faster version?
  - Adding more processors?

# Relative (Normalized) Performance

- Definition: Performance = 1/Execution Time
- “X is  $n$  times faster than Y”

$$\begin{aligned} \text{Performance}_X / \text{Performance}_Y \\ = \text{Execution time}_Y / \text{Execution time}_X = n \end{aligned}$$

- Example: time taken to run a program
  - 10s on A, 15s on B
  - $\text{Execution Time}_B / \text{Execution Time}_A$   
 $= 15\text{s} / 10\text{s} = 1.5$
  - So A is 1.5 times (or 50%) faster than B

# Measuring Execution Time

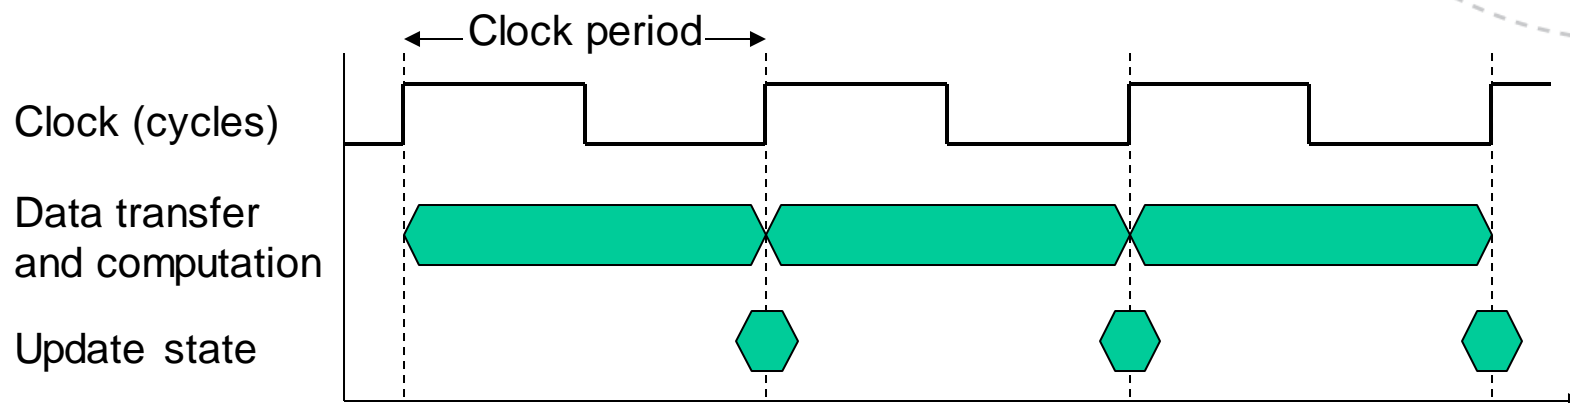
- Elapsed time/Wall clock time
  - Total turn-around time, including all aspects
    - Processing, I/O, OS overhead, idle time
  - Determines system performance
- CPU time
  - Time spent processing a given job
    - Discounts I/O time, other jobs' shares
  - Comprised of **user CPU time** and **system CPU time**
  - Different programs are affected differently by CPU and system performance
- Which programs to use?
  - Choose a representative set of real applications (i.e., benchmarks)

# Benchmarks

- Must use real applications
  - Programs tend to consist of phases
  - The relative importance of each phase creates bias – different computer architecture techniques pay off in different phases
  - Recreating this bias synthetically is very difficult
- Must use a collection of real applications
  - A computer is general-purpose – it must perform decently across as many applications as possible
- Benchmark suites:
  - SPEC CPU, SPEC OMP, PARSEC, NAS, TPC, etc.
  - Each suite attempts to be representative for the workloads of a chosen domain
- Discredited benchmarking approaches:
  - Kernels, toy programs and synthetic benchmark
  - Do not accurately recreate the behavior of a collection of real programs
  - Another problem: It is too easy to cheat if the benchmarks are simple.

# CPU Clocking

- Operation of digital hardware governed by a constant-rate clock



- Clock period: duration of a clock cycle
  - e.g.,  $250\text{ps} = 0.25\text{ns} = 250 \times 10^{-12}\text{s}$
- Clock frequency (rate): cycles per second
  - e.g.,  $4.0\text{GHz} = 4000\text{MHz} = 4.0 \times 10^9\text{Hz}$

# CPU Time

$$\begin{aligned}\text{CPU Time} &= \text{CPU Clock Cycles} \times \text{Clock Cycle Time} \\ &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}}\end{aligned}$$

- Performance improved by
  - Reducing number of clock cycles
  - Increasing clock rate
- Hardware designer must often trade off clock rate against cycle count

# CPU Time Example

- Computer A: 2GHz clock, 10s CPU time on benchmark X
- Designing Computer B
  - Design Goal: 6s CPU time on benchmark X
  - Faster clock causes 20% increase in the number of clock cycles
- How fast must Computer B clock be?

$$\text{Clock Rate}_B = \frac{\text{Clock Cycles}_B}{\text{CPU Time}_B} = \frac{1.2 \times \text{Clock Cycles}_A}{6s}$$

$$\begin{aligned}\text{Clock Cycles}_A &= \text{CPU Time}_A \times \text{Clock Rate}_A \\ &= 10s \times 2\text{GHz} = 20 \times 10^9\end{aligned}$$

$$\text{Clock Rate}_B = \frac{1.2 \times 20 \times 10^9}{6s} = \frac{24 \times 10^9}{6s} = 4\text{GHz}$$

# Instruction Count and CPI

$\text{Clock Cycles} = \text{Instruction Count} \times \text{Cycles per Instruction}$

$\text{CPU Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$

$$= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$$

- Instruction Count for a program
  - Determined by program, ISA and compiler
- Average Cycles Per Instruction (CPI)
  - Determined by CPU hardware
  - If different instructions have different CPI
    - Average CPI affected by instruction mix
  - Sometimes we report the inverse of CPI – called Instructions Per Cycle (IPC) since this is a higher-is-better metric
- Under which conditions is it valid to compare CPIs directly?



# CPI Example

- Computer A: Cycle Time = 250ps, CPI = 2.0
- Computer B: Cycle Time = 500ps, CPI = 1.2
- Same ISA, program and compiler
- Which is faster, and by how much?

$$\begin{aligned}\text{CPU Time}_A &= \text{Instruction Count} \times \text{CPI}_A \times \text{Cycle Time}_A \\ &= 1 \times 2.0 \times 250\text{ps} = 1 \times 500\text{ps} \quad \leftarrow \text{A is faster...}\end{aligned}$$

$$\begin{aligned}\text{CPU Time}_B &= \text{Instruction Count} \times \text{CPI}_B \times \text{Cycle Time}_B \\ &= 1 \times 1.2 \times 500\text{ps} = 1 \times 600\text{ps}\end{aligned}$$

$$\frac{\text{CPU Time}_B}{\text{CPU Time}_A} = \frac{1 \times 600\text{ps}}{1 \times 500\text{ps}} = 1.2 \quad \leftarrow \text{...by this much}$$

# CPI in More Detail

- If different instruction classes take different numbers of cycles

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

- Weighted average CPI

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left( \text{CPI}_i \times \frac{\text{Instruction Count}_i}{\text{Instruction Count}} \right)$$

Relative frequency

# CPI Example

- Alternative compiled code sequences using instructions in classes A, B, C

Class	A	B	C
CPI for class	1	2	3
IC in sequence 1	2	1	2
IC in sequence 2	4	1	1

## ■ Sequence 1: IC = 5

- Clock Cycles  
 $= 2 \times 1 + 1 \times 2 + 2 \times 3$   
 $= 10$
- Avg. CPI =  $10/5 = 2.0$

## ■ Sequence 2: IC = 6

- Clock Cycles  
 $= 4 \times 1 + 1 \times 2 + 1 \times 3$   
 $= 9$
- Avg. CPI =  $9/6 = 1.5$

# The Iron Law

$$\text{CPU Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$

- Performance depends on
  - Algorithm: affects IC, possibly CPI
  - Programming language: affects IC, CPI
  - Compiler: affects IC, CPI
  - Instruction set architecture: affects IC, CPI, Clock cycle time

# SUMMARY

# Summary

- CPUs are pervasive and ubiquitous
  - Fully understanding how they are constructed is a good preparation for a career in ICT
- Performance of single-core processors have saturated
  - Future processors are likely to be heterogeneous
  - Understanding the basics of computer design is very helpful when trying to use these systems efficiently
- Key metric: Performance
  - Normally traded off against chip area, energy, power, etc.
  - Measured on a collection of programs (benchmarks) that are (hopefully) representative of a given domain.