



# NTNU

Det skapende universitet

**TDT4255 Computer Design**

**Lecture 5: Branch Prediction**

**Magnus Jahre**

# Outline

- Appendix C.2
- Chapter 3.3

# BRANCH PREDICTION

*Slides in this section are by Lieven Eeckhout, Ghent University.  
Reused with permission.*

# Instruction fetch

- Goal is to fetch as many insns per cycle as possible
- Branch prediction is super important to fill the pipeline with correct-path instructions (useful work)
- Importance increases with deeper pipelines

# Branch prediction

- Goal is to predict the branch direction and target address, and start fetching and executing insns along the predicted path
- Key observation:
  - Branches exhibit temporal locality
  - Predicting branch behavior
    - Keep track of past history
    - Predict the future based on the past
  - Branch behavior is predictable
    - typically over 90%, 95% or 99% of all dynamically executed branches are correctly predicted

# Cost per mispredicted branch

- Cost per mispredicted branch is proportional to pipeline depth (no. pipeline stages)
- Early pipelines had 4-6 stages
- Modern pipelines have 12-15 stages
  - Example: Intel Pentium 4 had 20 stages (Extreme Edition had 31 stages)
  - Reason: Higher clock frequencies

# Branch Direction Prediction

- Applies to conditional branches
- Static prediction
  - Static = before program execution
  - One prediction (taken/not-taken) per static branch in the program binary
  - Via software: compiler or programmer
- Dynamic prediction
  - Dynamic = during program execution
  - Multiple predictions per static branch, depending on history (= outcomes of prior branch executions) of that particular branch or even other branches
  - Done in hardware

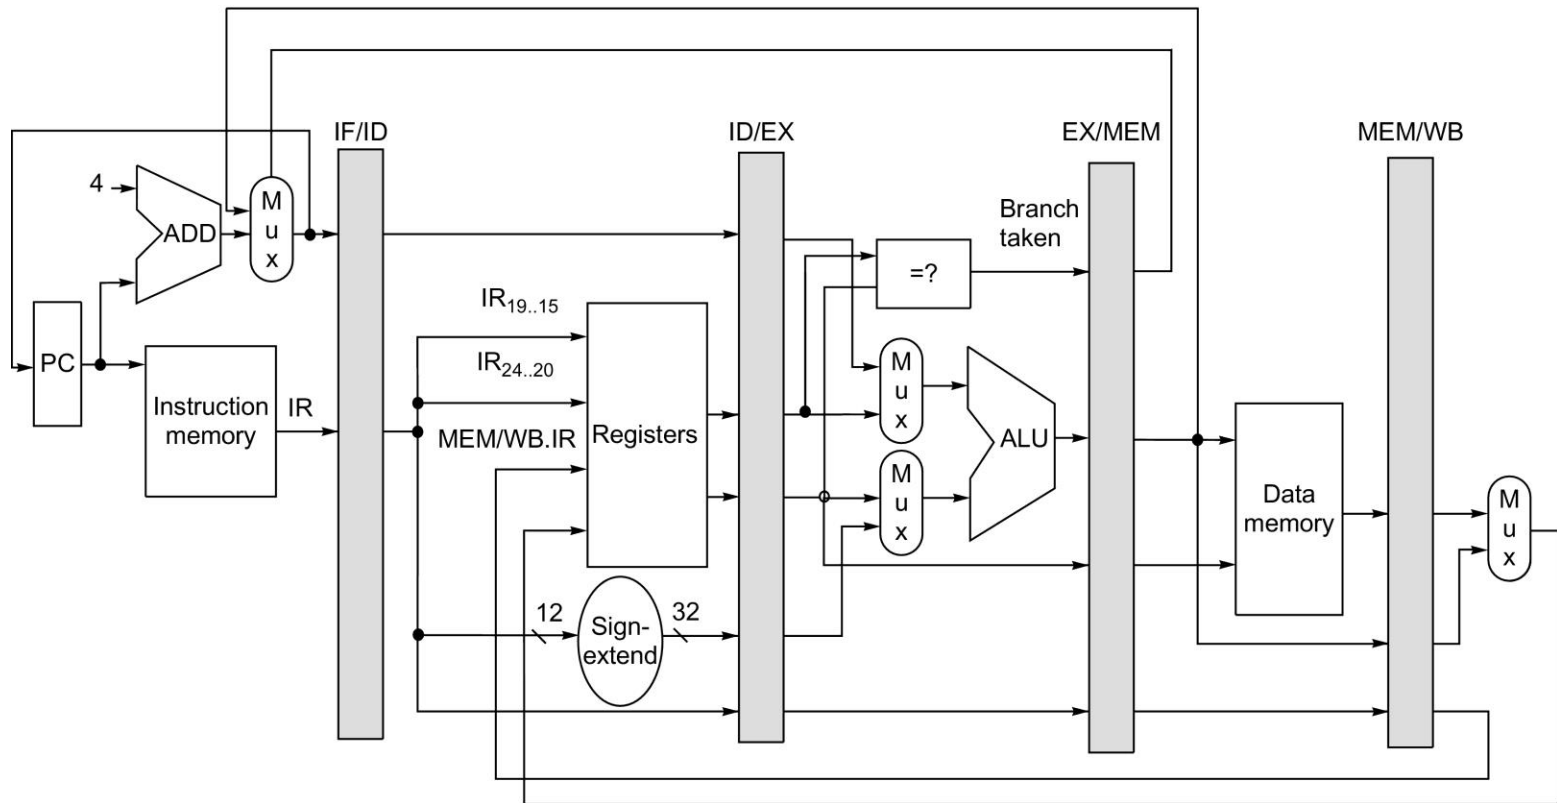
# Branch address vs. branch target address

```
000000012002e530 <__start>:
 12002e530: f0 ff de 23      lda     sp, -16(sp)
 12002e534: 08 00 fe b7      stq     zero, 8(sp)
 12002e538: 00 00 20 c0      br      t0, 12002e53c <__start+0xc>
 12002e53c: 10 00 1e a2      ld1     a0, 16(sp)
 12002e540: 18 00 3e 22      lda     a1, 24(sp)
 12002e544: 00 20 a1 27      ldah    gp, 8192(t0)
 12002e548: 52 06 11 42      s8addq  a0, a1, a2
 12002e54c: 12 14 41 42      addq    a2, 0x8, a2
 12002e550: 24 28 bd 23      lda     gp, 10276(gp)
 12002e554: 13 04 52 46      mov     a2, a3
 12002e558: 00 00 33 a4      ldq     t0, 0(a3)
 12002e55c: 13 14 61 42      addq    a3, 0x8, a3
 12002e560: fd ff 3f f4      bne     t0, 12002e558 <__start+0x28>
```

Command on Linux/UNIX/OS X: `objdump -d a.out`

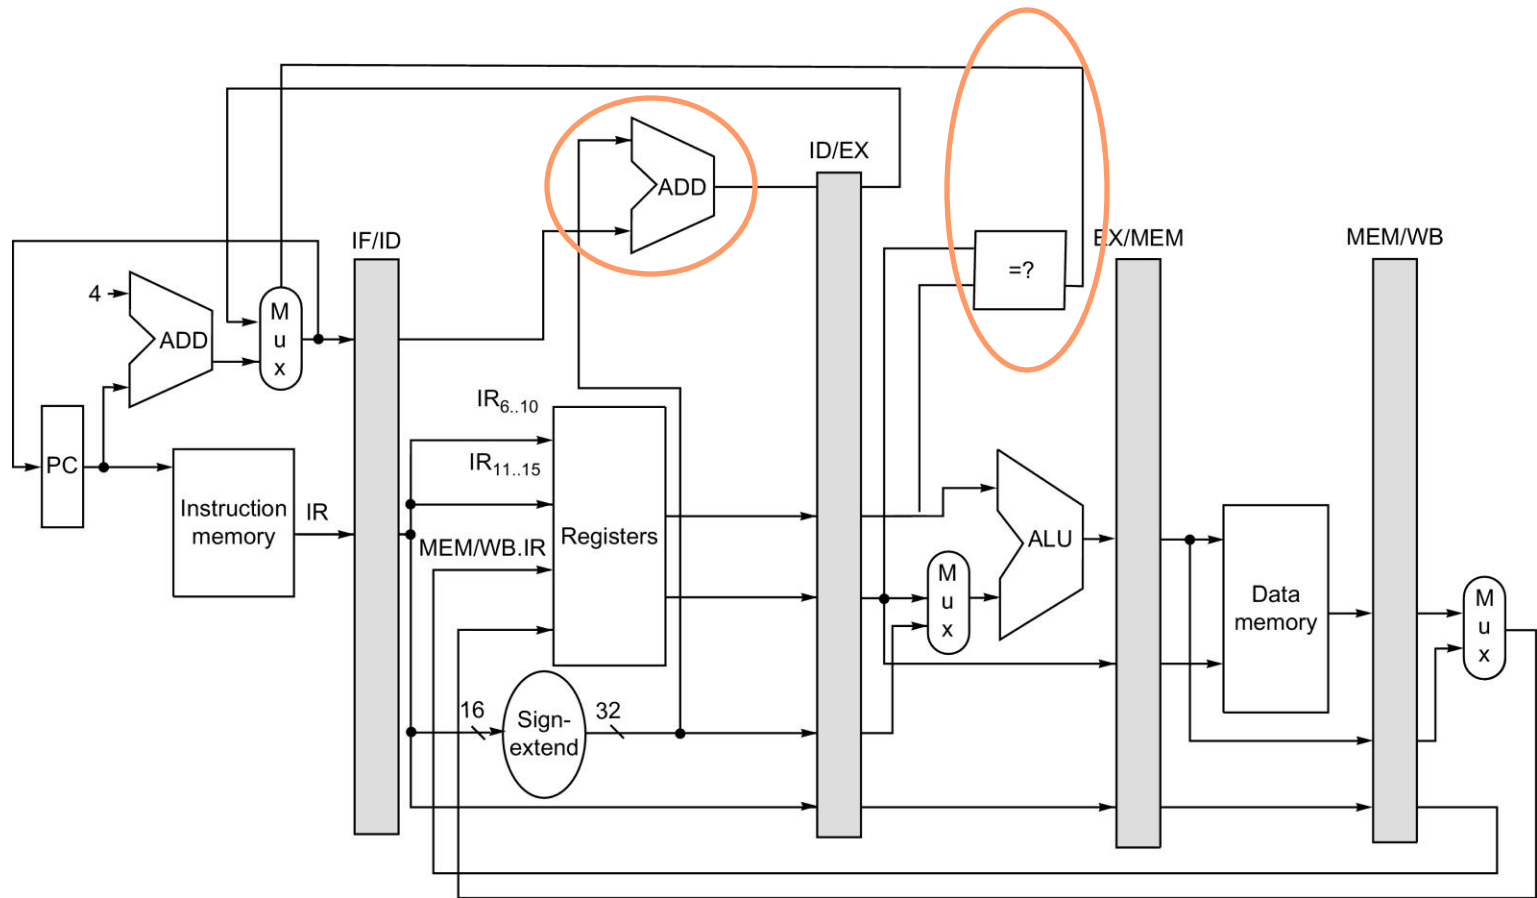


# Vanilla Pipeline



Branch penalty: Three cycles

# Faster Conditional Branches



Branch penalty: Two cycles

# STATIC BRANCH DIRECTION PREDICTION

*Slides in this section are by Lieven Eeckhout, Ghent University.  
Reused with permission.*

# Static Branch Prediction

- Advantages
  - Easy to implement
  - Little HW is needed
- Disadvantages
  - Provides the same prediction regardless of input and/or dynamic execution behavior
- Three flavors
  - rule-based
  - program-based
  - profile-based

# Rule-based

- Always not-taken
  - Simple HW: sequential fetch
- Always taken
  - HW is more complex because of unknown branch target
  - Branch target is known at the decode stage
  - May lead to a *bubble* ( ‘lost cycle’ ) in the pipeline
- BTFNT
  - *Backward taken - forward not-taken*
  - Branches to smaller addresses are typically backward loop branches and are typically taken

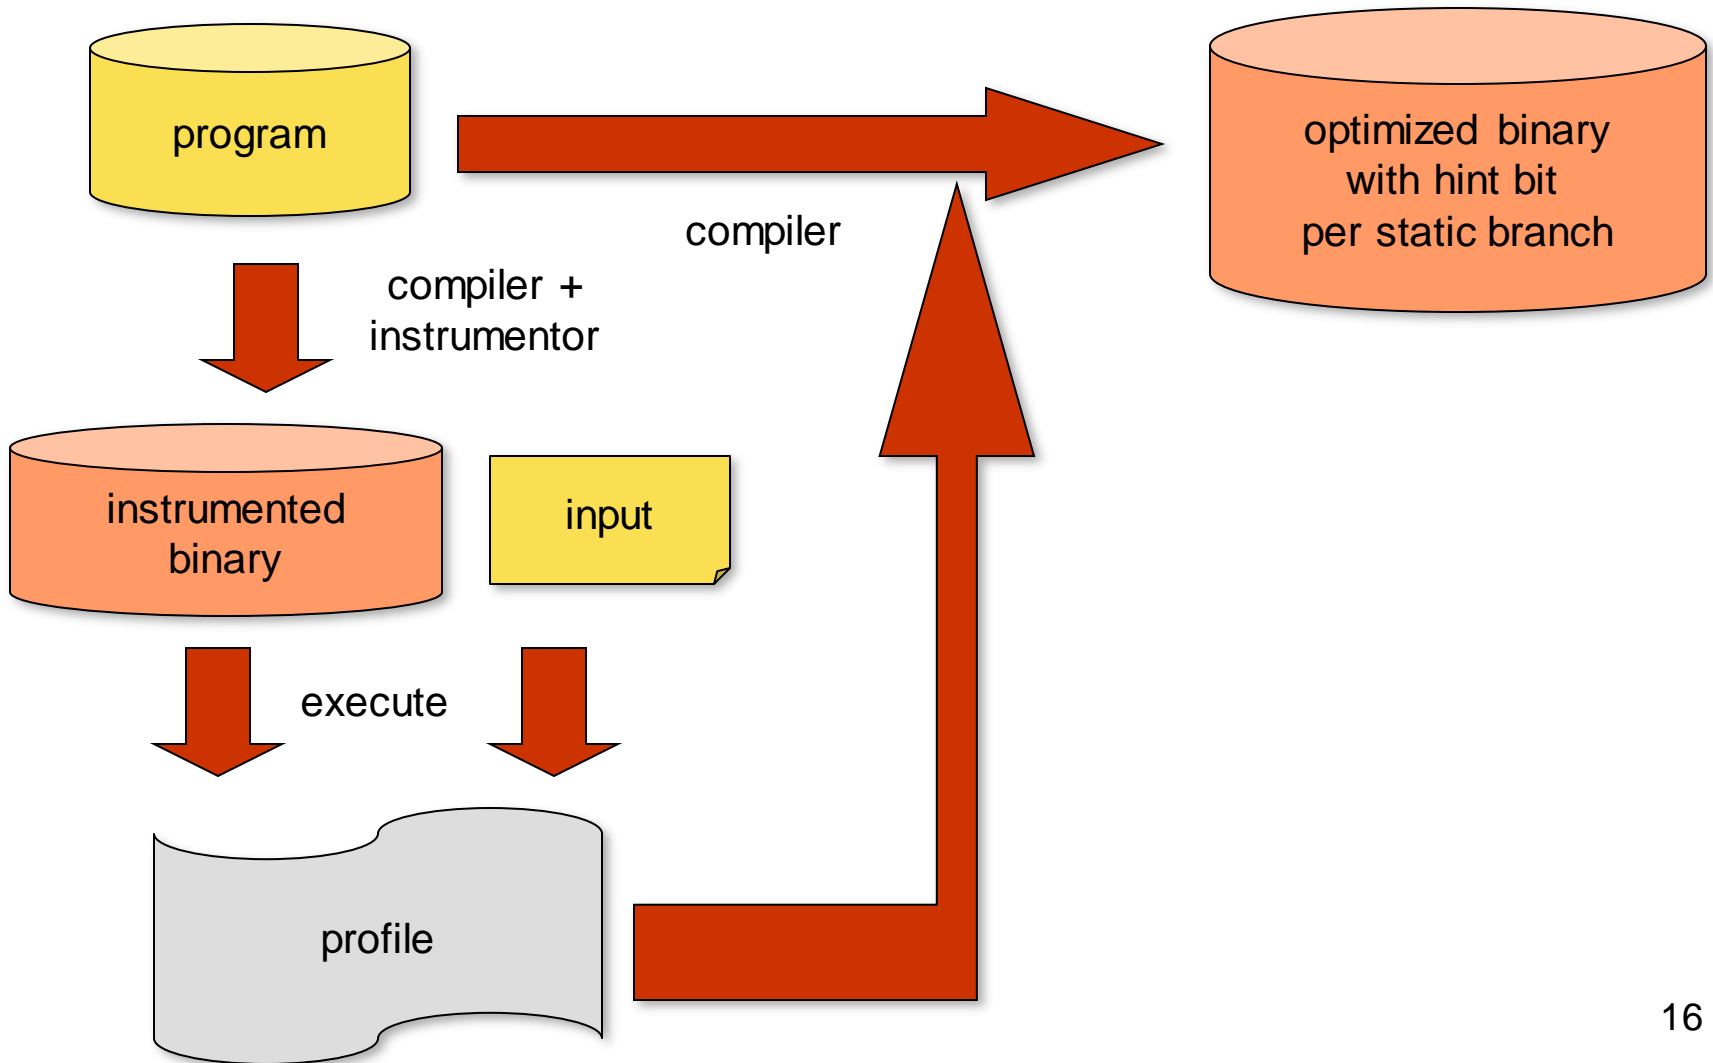
# Program-based

- Ball and Larus heuristics (PLDI'93)
  - Requires a hint bit in instruction opcode
  - Branch direction is estimated based on program structure
  - Examples:
    - Predict loop branches to be taken
    - When comparing a pointer to NULL, predict branch direction to non-NULL path
    - When comparing two pointers, predict branch direction to path representing pointer inequality
  - Typically more accurate than rule-based

# Profile-based

- Execute instrumented binary with a given training input to collect profile information
  - Count how often a static branch is taken/not-taken
- Use profile information during recompilation and add hint bits
  - Predict taken if branch has higher probability than 50% to be taken; predict not-taken otherwise
- Typically more accurate than rule- and program based


# Profile-based





# Static Branch Prediction Accuracy

- Some branches are biased towards being taken and some are biased towards being not taken
- A single, static policy for all branches is unlikely to work
- Even using profiling to determine the preferred direction of a branch is not sufficient (e.g., alternating between taken and not taken)



# DYNAMIC BRANCH DIRECTION PREDICTION

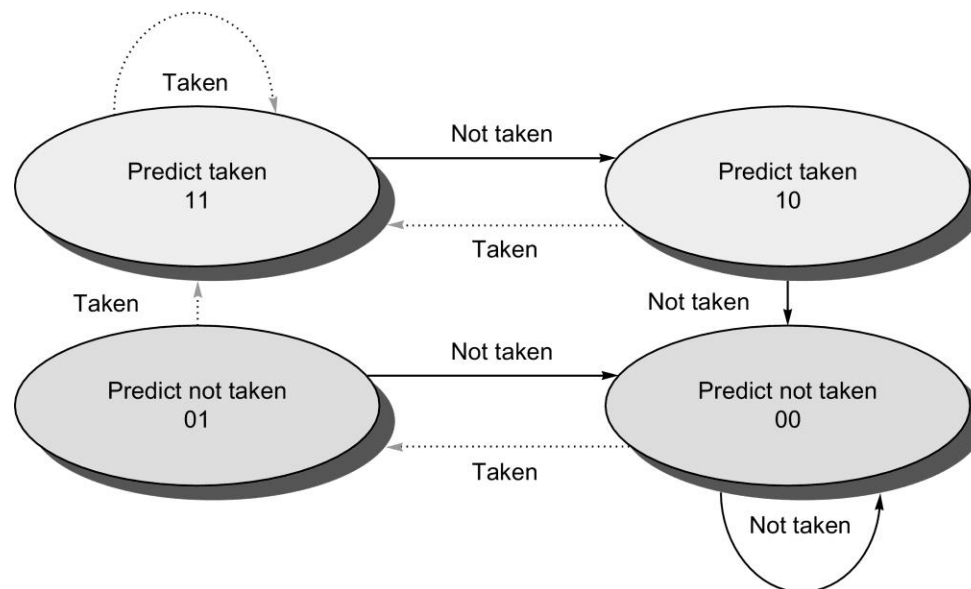
*Slides in this section are by Lieven Eeckhout, Ghent University.  
Reused with permission.*

# Dynamic branch prediction

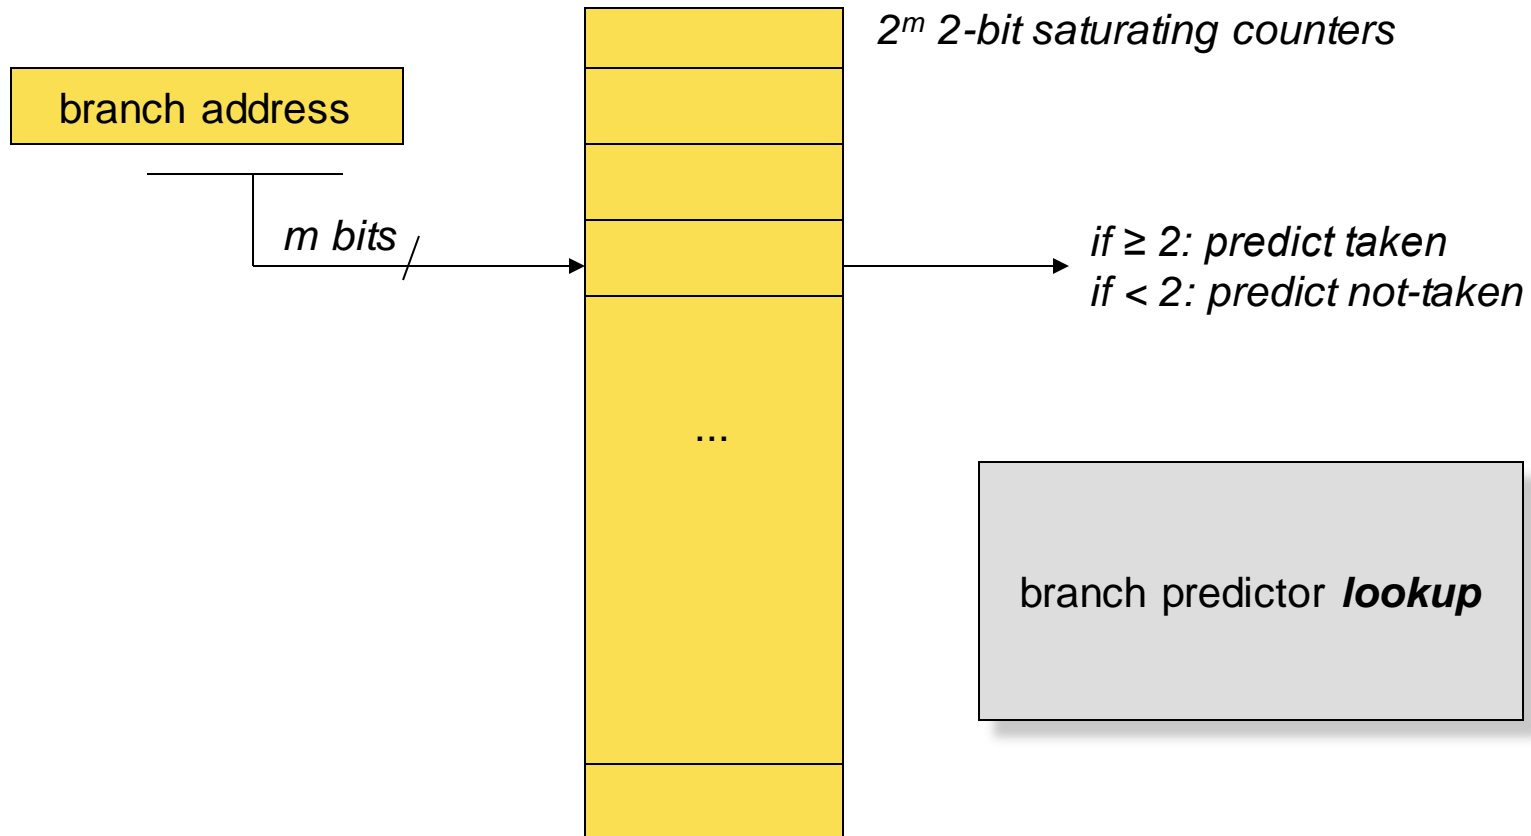
- More accurate than static branch prediction
  - 80%-97% (dynamic) vs. 50%-80% (static)
- Some branches are hard to predict statically, but are easily predicted dynamically
  - Some examples
    - Not-taken during first half of execution, and taken during second half
    - Alternating taken/not-taken
- Takes into account branch context!
  - the branch's own history (***local history***)
  - other branches' histories (***global history***)

# Bimodal predictor

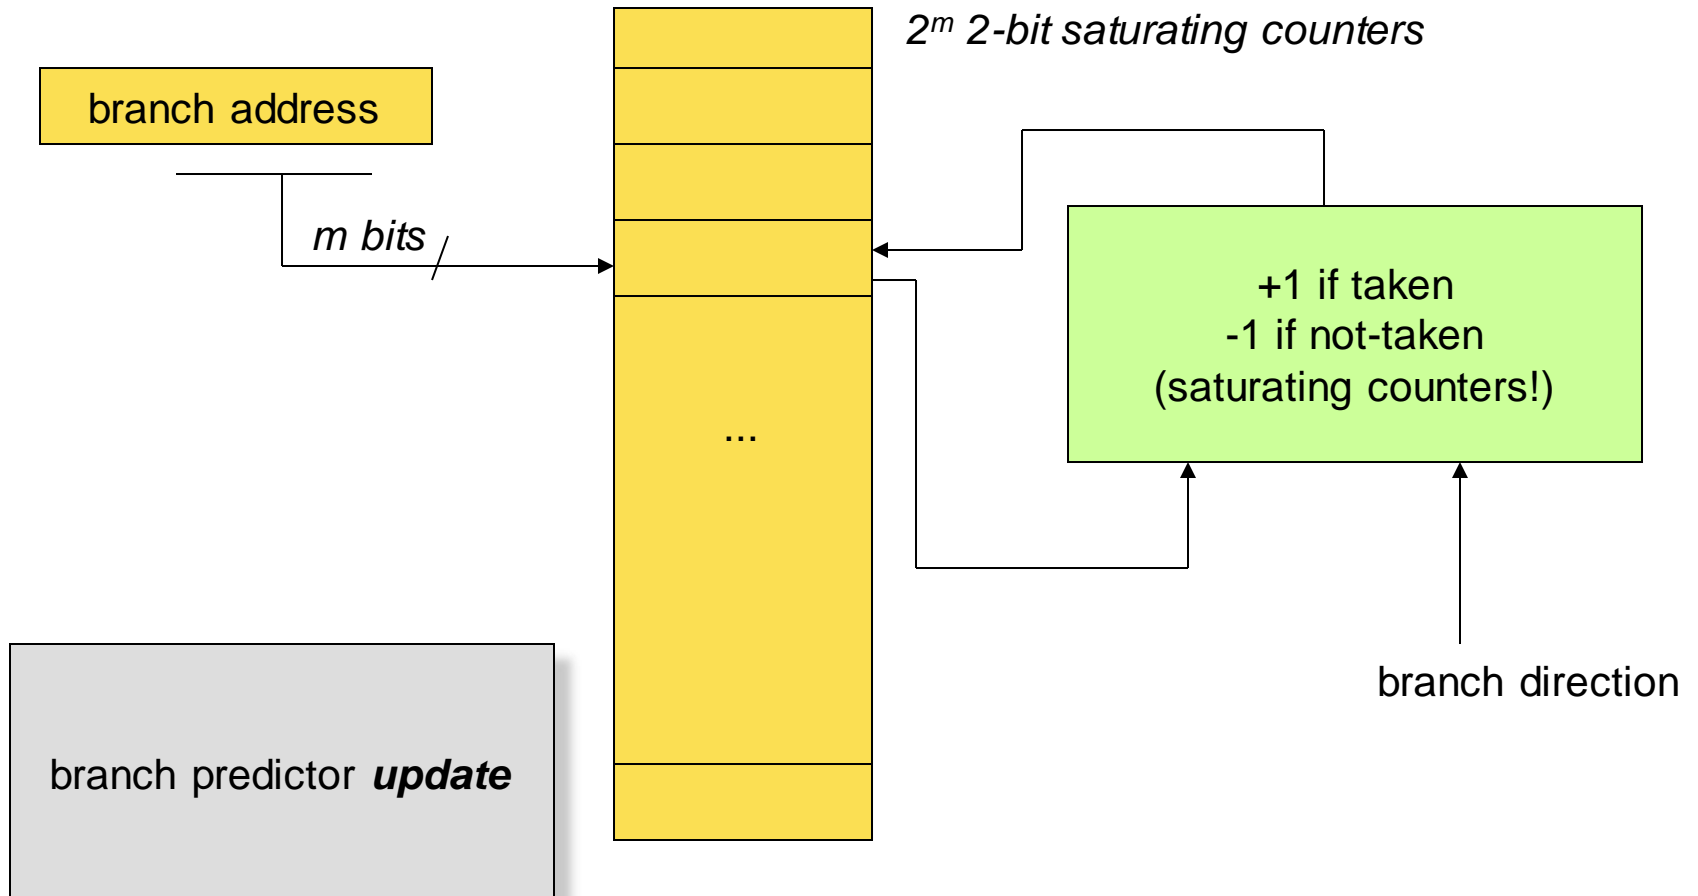
- Hash the PC of a branch into a small buffer which stores the branch outcome and use this for making predictions
- A 2-bit predictor is does not change its prediction when it is wrong once (which typical occur on branches that check loop conditions):



# Bimodal predictor



# Bimodal predictor




# Example bimodal predictor

<i>old state</i>	<i>prediction</i>	<i>branch direction</i>	<i>new state</i>
00	0	0	00
00	0	0	00
00	0	1	01
01	0	1	10
10	1	1	11
11	1	1	11
11	1	0	10
10	1	1	11
11	1	1	11
11	1	0	10
10	1	0	01

*Predictions and predictor state for a single PC (i.e., single branch) over time*

# How good is the 2-bit predictor?

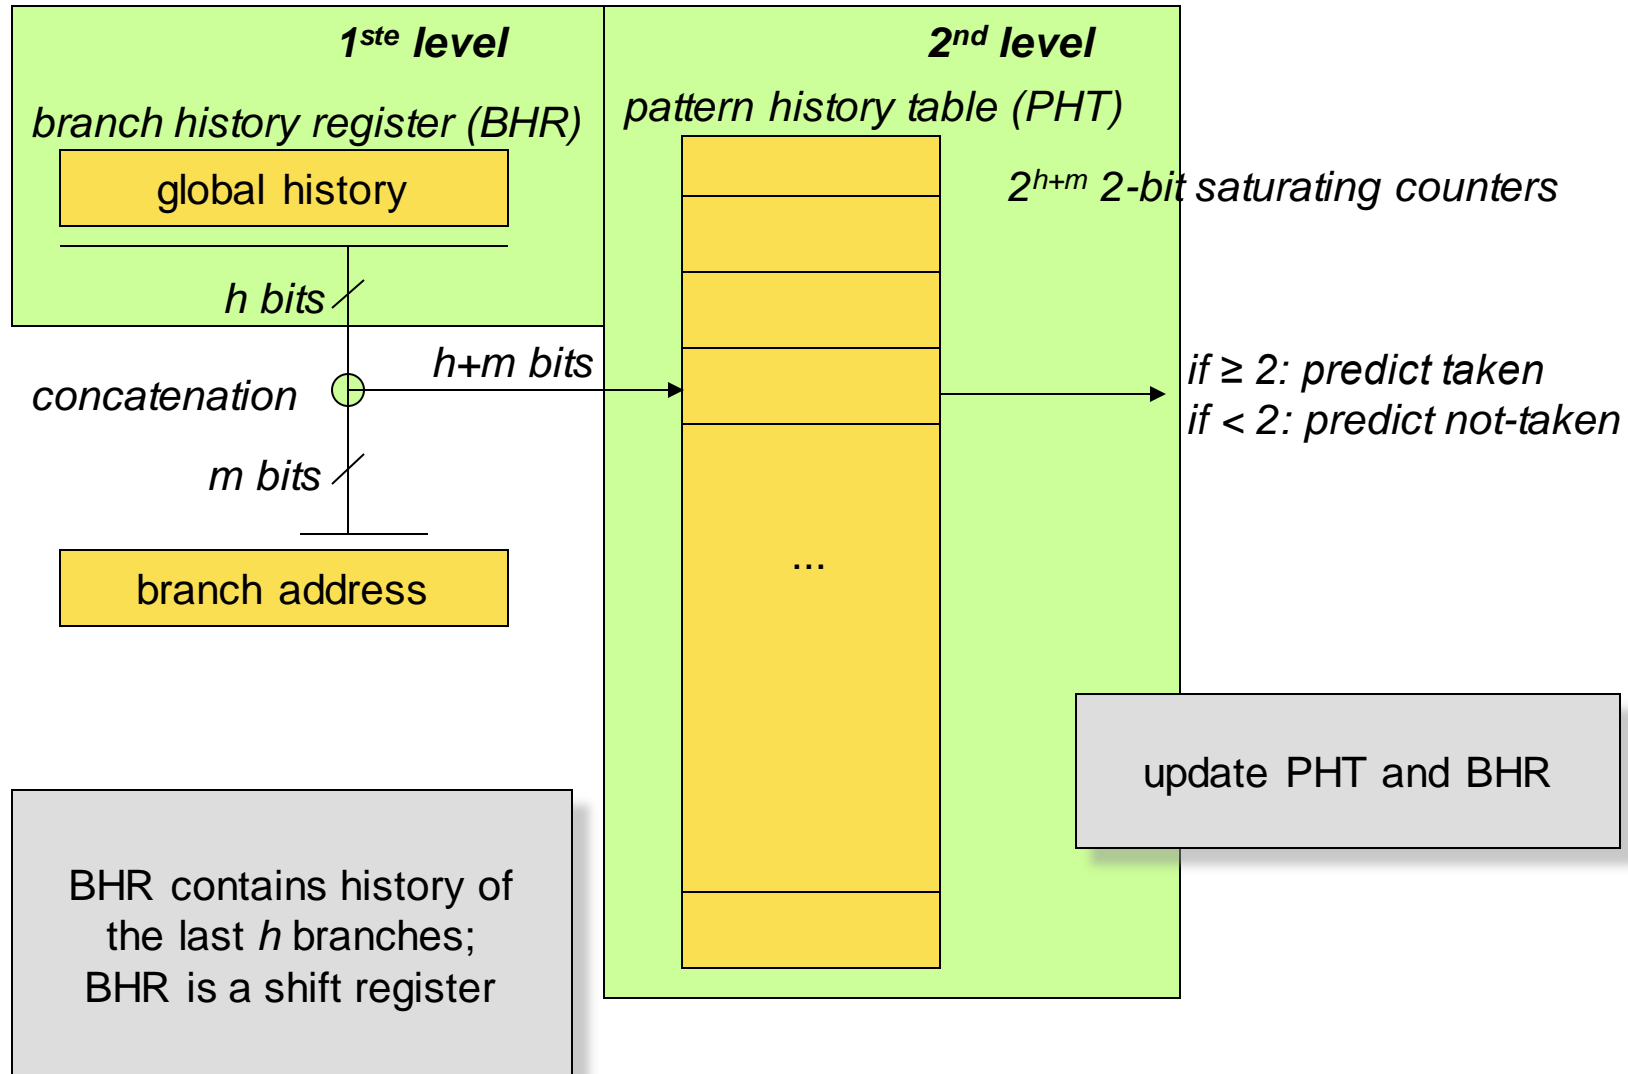




# Two-level predictors

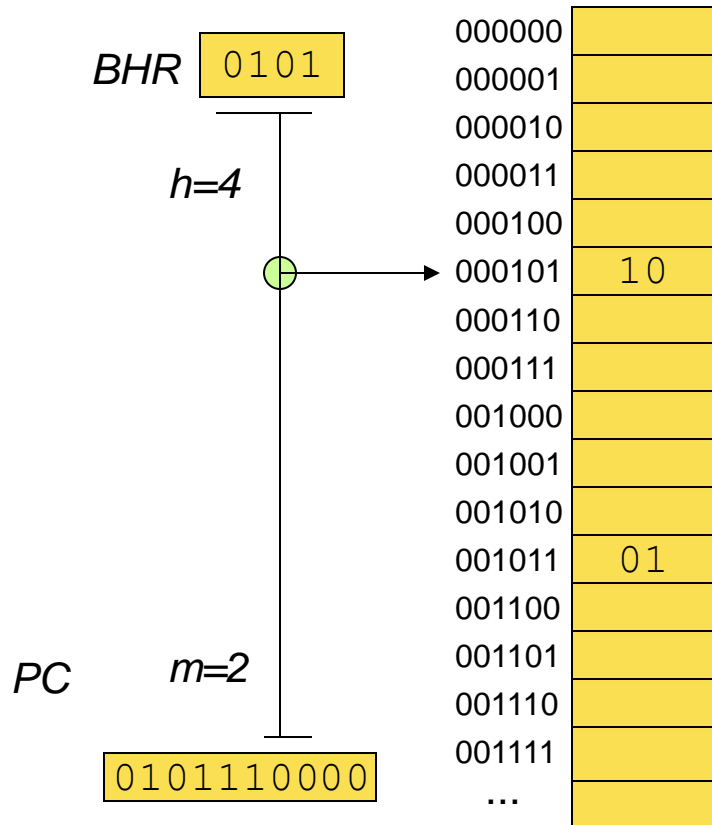
- *Correlating branch predictor*
- Uses (two levels of) branch history to make a prediction
  - Bimodal predictor uses the branch address only
  - Two-level predictors also use
    - *Local history* of that particular branch, or
    - *Global history* of all prior branches
    - history = outcome of prior branch executions
    - history acts as a shift register: newest branch outcome is inserted; oldest branch outcome is discarded
- Proposed by Tse-Yu Yeh and Yale N. Patt in 1991

# ... using global history




# Example

Before



After



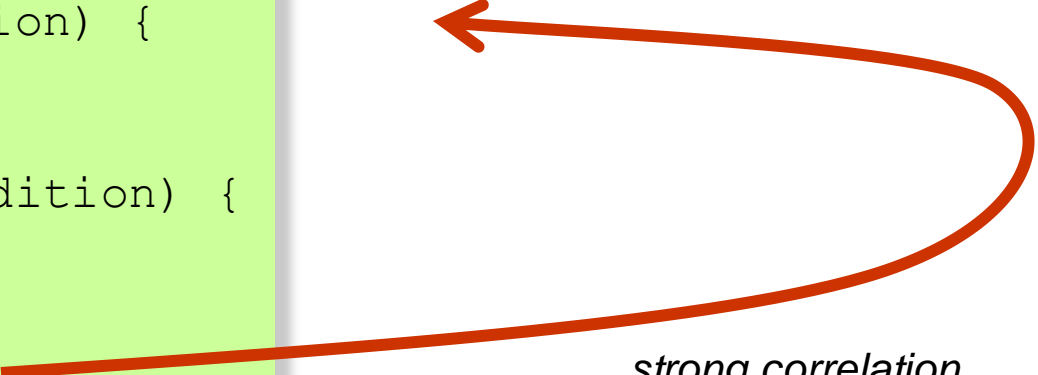
Assume: branch direction is taken

# Why does this work?

- Correlation among branches
  - For example, if a particular branch is taken (not-taken), probability may be high that next branches are also taken (not-taken), or vice versa
- Concrete examples
  - Branch conditions that depend on the same variable
  - A variable is control dependent on branch; a subsequent branch is data dependent on the variable; both branches will be correlated

# Example

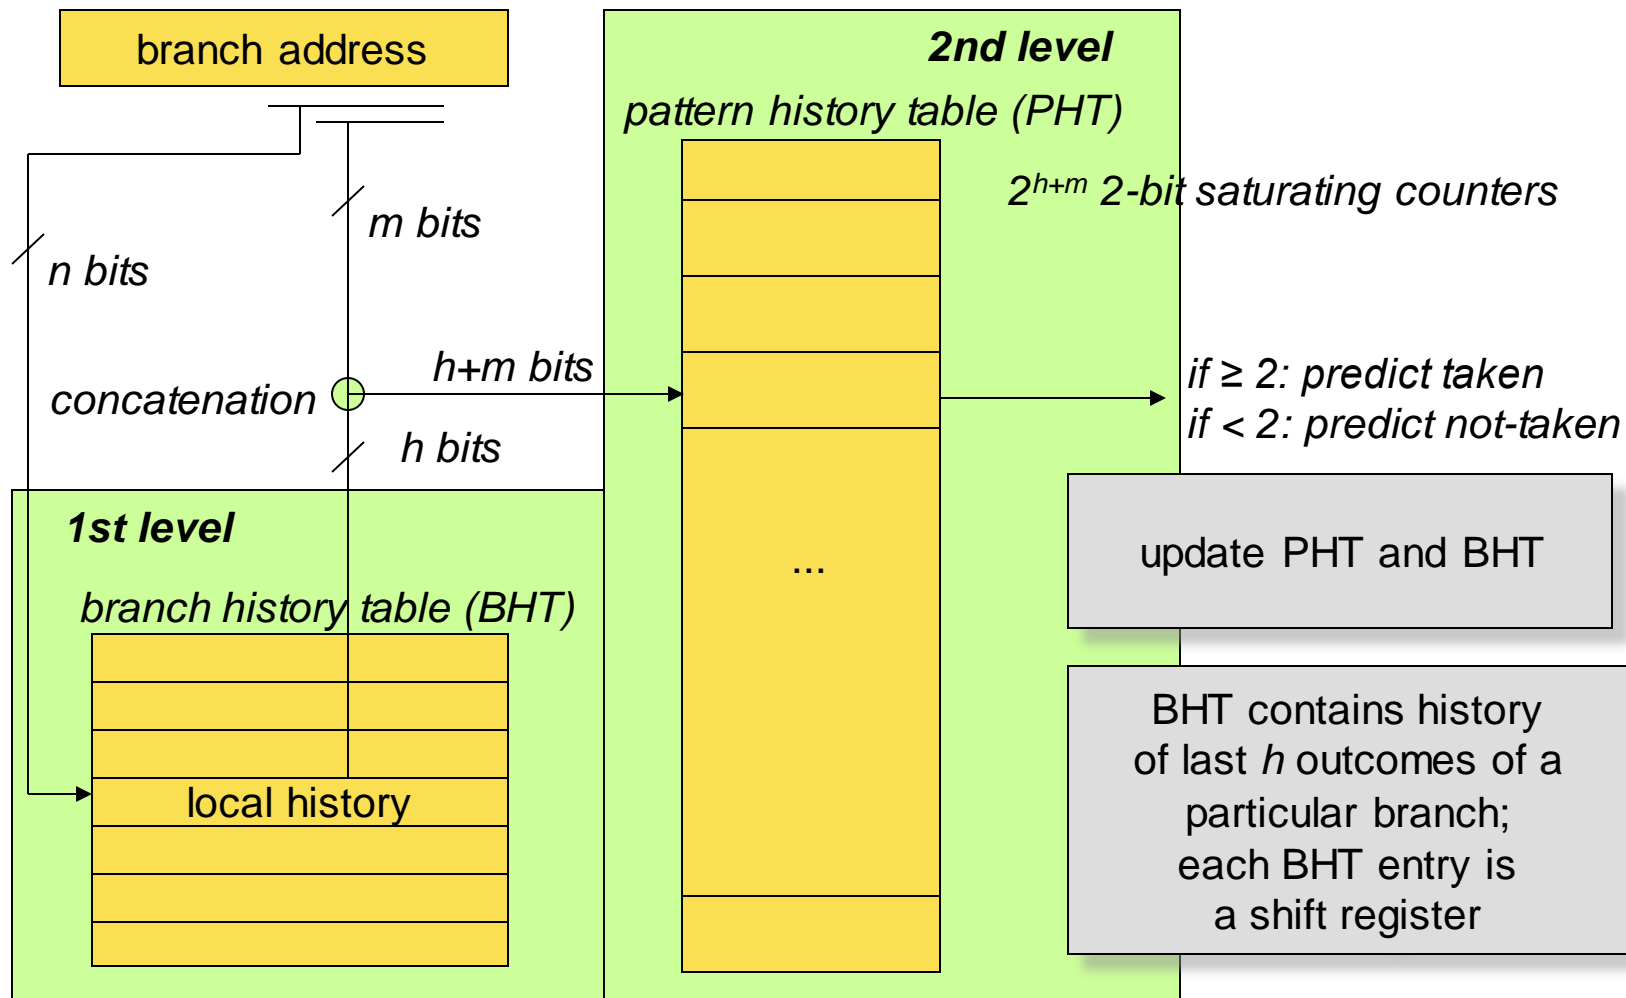
```
x = 0;  
if (some_condition) {  
    x = 3;  
}  
if (another_condition) {  
    y += 19;  
}  
if (x <= 0) {  
    do_something ();  
}
```



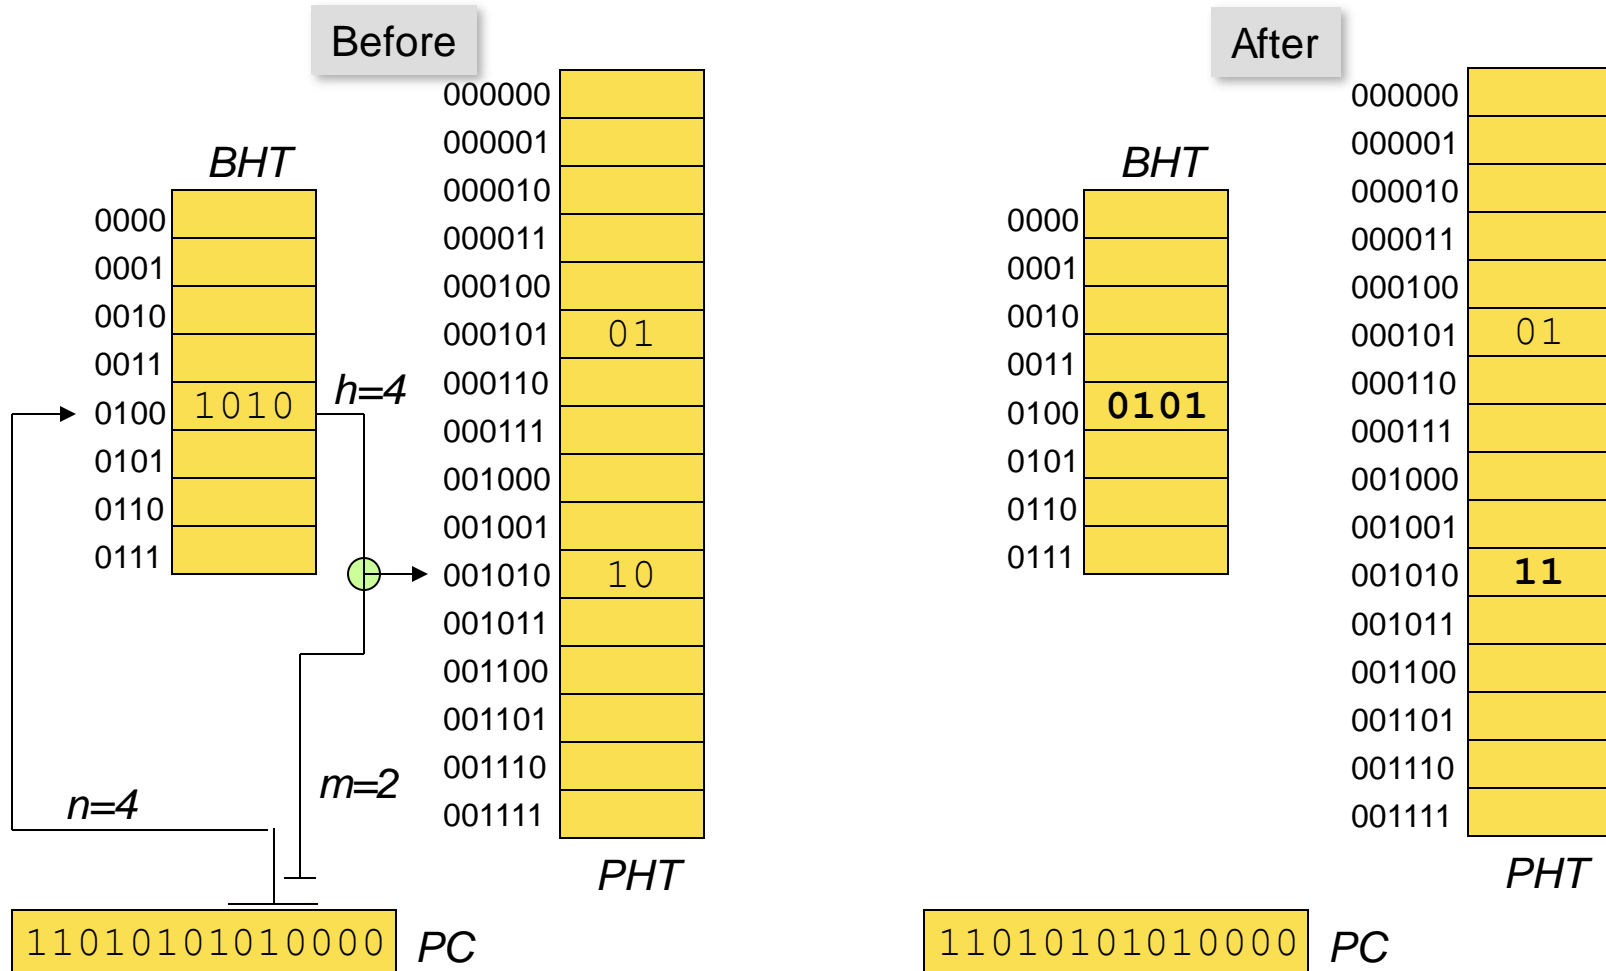
*strong correlation  
between 1st and 3rd branch*

Note: there is no correlation between  
the 2nd and 3rd branch → branch  
prediction needs to learn!

# ... using local history



# Example



Assume: branch direction is taken

# Why does this work?

- Local behavior of a branch
- Examples
  - Alternating: 01010101...
  - Loops with limited number of iterations:  
111011101110...



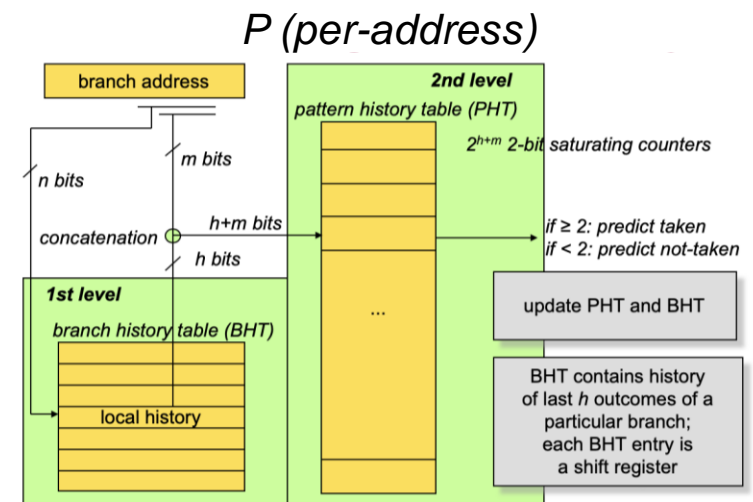
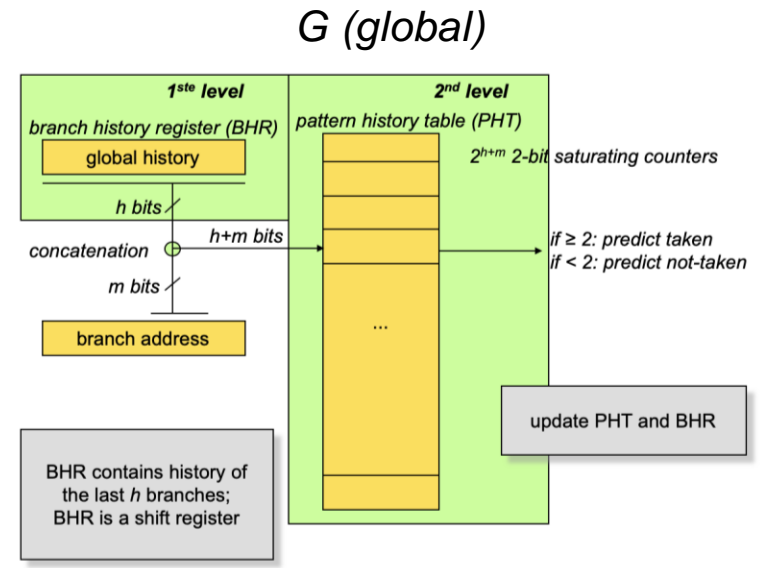
# Two-level predictor implementations

- If  $m=0$ 
  - *Global pattern history table (gPHT)*
  - All branches share the same PHT
- If  $m \neq 0$ 
  - *Per-address pattern history table (pPHT)*
  - PHT is partitioned based on branch address bits
- Four variations: GAg, GAp, PAg, PAp
  - ‘G’ global history, ‘P’ per-address (local)
  - ‘A’ adaptive
  - ‘g’  $\rightarrow$  gPHT and ‘p’  $\rightarrow$  pPHT

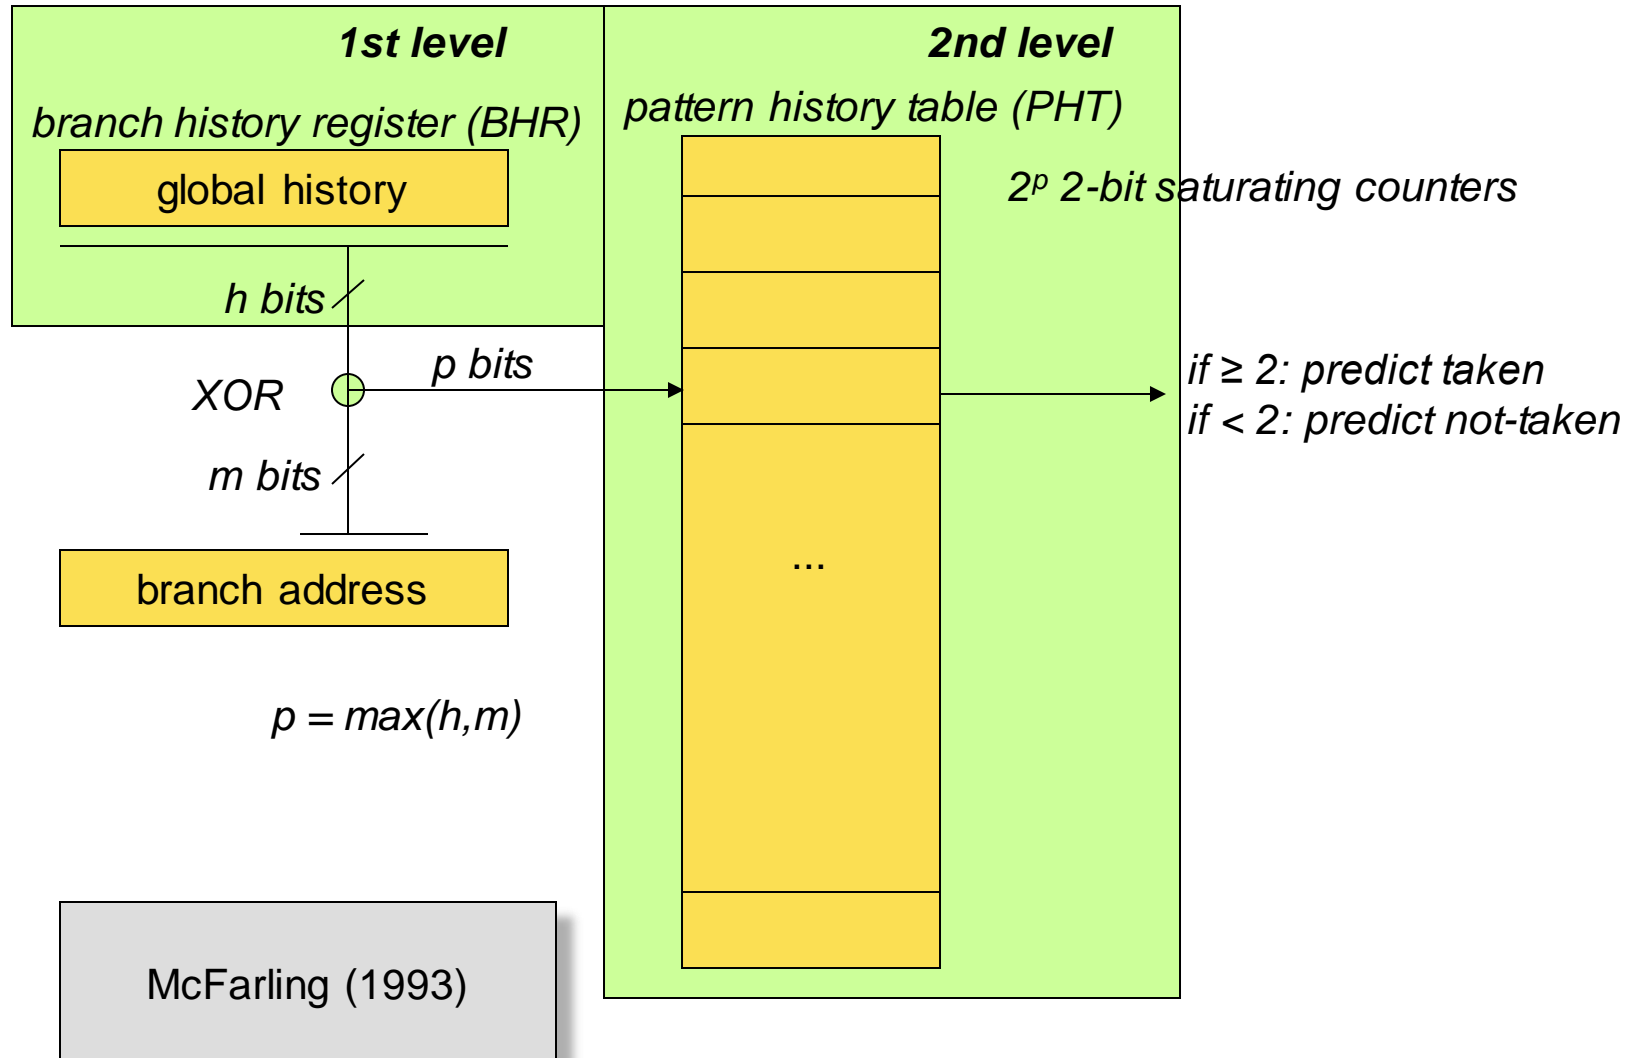
# Good configurations

- Obtained through experimental evaluation by Yeh/Patt (1991)
- GAg
  - BHR: 18 bits
  - PHT:  $2^{18} \times 2$  bits
- PAg
  - BHT:  $2^{11} \times 12$  bits
  - PHT:  $2^{12} \times 2$  bits
- PAp
  - BHT:  $2^{11} \times 6$  bits
  - PHT:  $2^9 \times 2^6 \times 2$  bits
- Prediction accuracy of 97%

*Remember: gPHT:  $m=0$*



# gshare



# GAp vs. gshare

- GAp needs to make a choice
  - How many bits to concatenate from BHR vs. branch address?
  - e.g., 5 bits from BHR and 5 bits from branch address
- gshare doesn't need to make this choice
  - e.g., 10 bits from BHR and 10 bits from branch address
  - More information is used to index the PHT

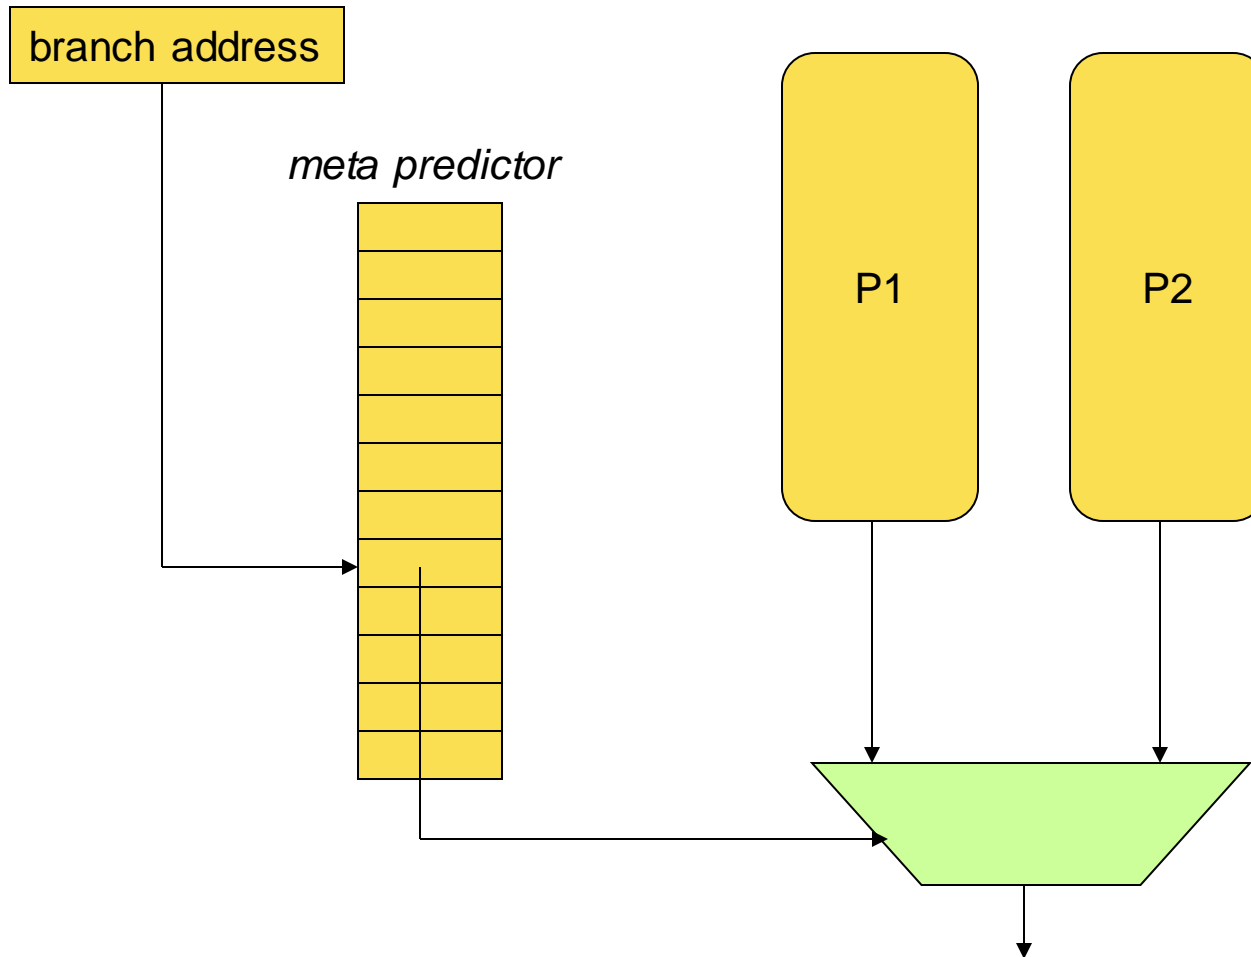
# Towards hybrid branch prediction

- Reasons for mispredictions?
- Branches may be hard to predict
  - During training phase of the branch predictor, e.g., after a context switch
- Interference or aliasing
  - Mostly negative interference
    - Branch predictor is too small for the number of branches
    - Hence, branches with different behavior may be mapped onto the same PHT entry
- Branch behavior doesn't match branch predictor type

# Hybrid branch prediction

- Idea: Combine different types of branch predictors and learn which one is most accurate for which branch

# Tournament predictor



# Tournament predictor

- Meta predictor determines which predictor will be followed
  - If  $< 2 \rightarrow P1$ ; If  $\geq 2 \rightarrow P2$
- Update meta predictor
  - If both correct or incorrect: do nothing
  - If P1 is correct and P2 incorrect: decrement
  - If P1 is incorrect and P2 is correct: increment
- Both predictors are updated
  - Irrespective of whether P1 or P2 was followed



# Tournament predictor (cont.)

- Various flavors exist
- Any predictor can be part of a hybrid predictor
- Indexing of meta predictor is to be chosen
- Hybrid predictor typically combines local and global history based predictors

# Alpha 21264

- Hybrid predictor
  - PAg
    - 1st level: 1K 10-bit entries
    - 2<sup>nd</sup> level: 1K 3-bit saturating counters
  - GAg
    - 4K 2-bit saturating counters
    - 12-bit global history
  - Meta predictor
    - 4K 2-bit saturating counters
    - Indexed using 12-bit global history

# IBM POWER4

- Hybrid predictor
  - Bimodal predictor
    - 16K 1-bit saturating counters
  - Gshare
    - 16K 1-bit saturating counters
    - 11-bit global history
  - Meta predictor
    - 16K 1-bit saturating counters
    - gshare indexing

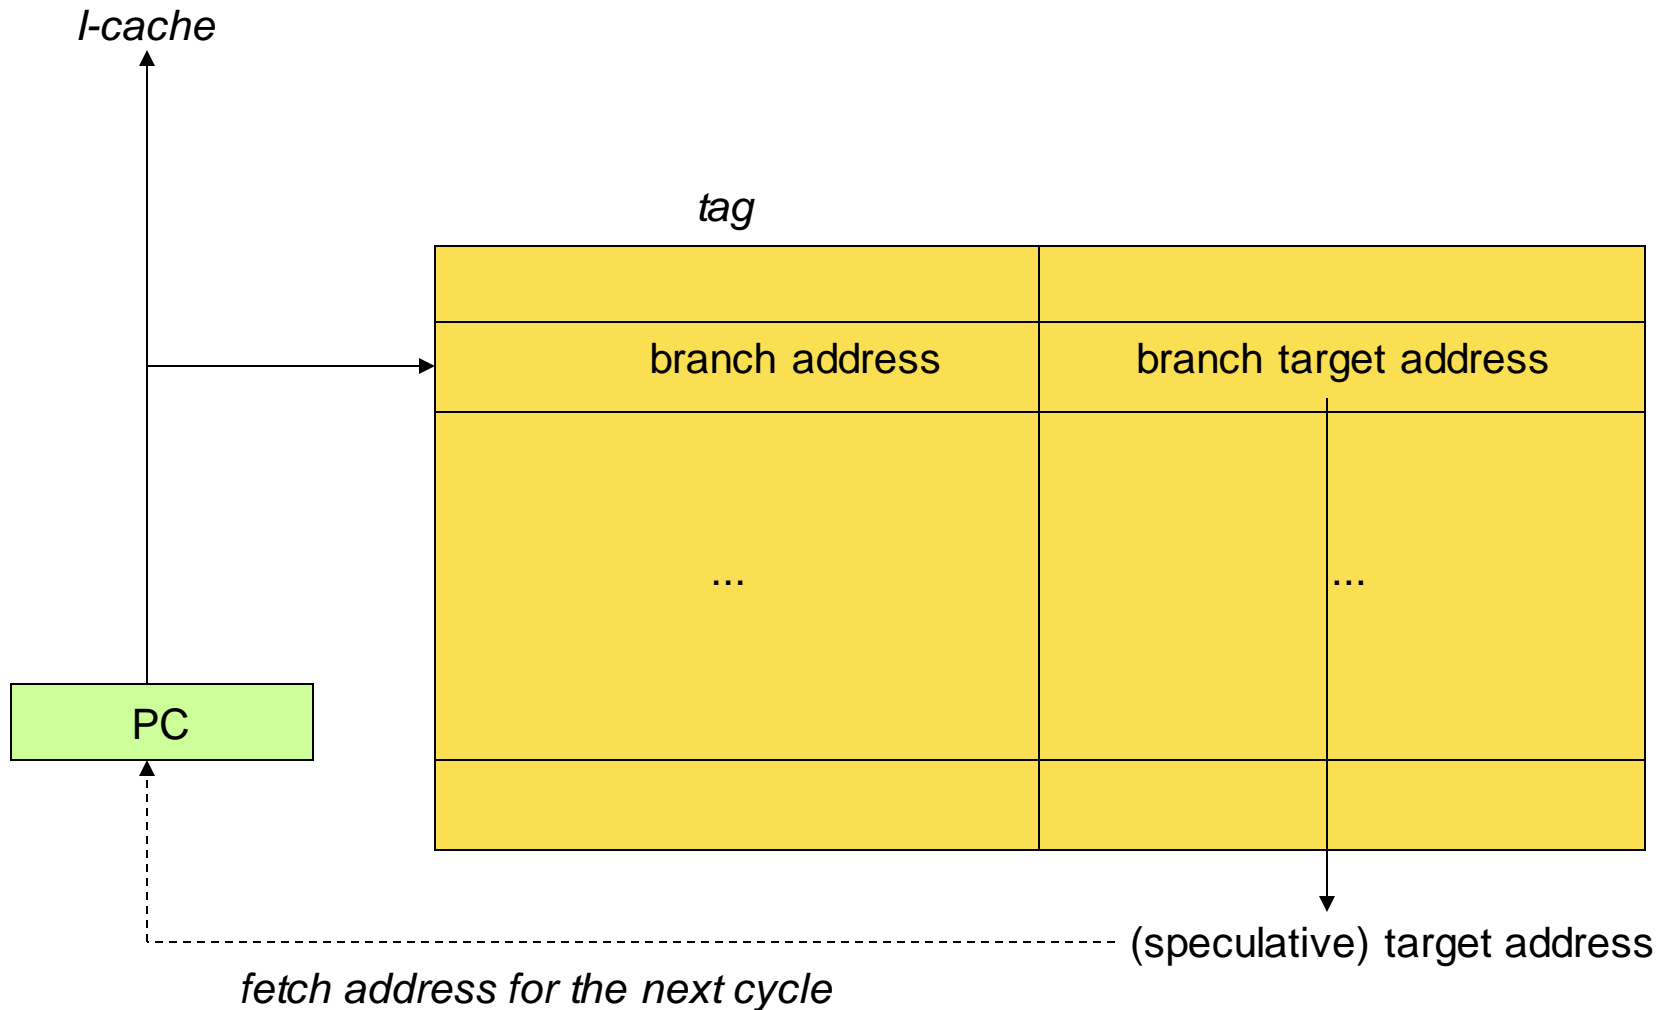
# BRANCH TARGET PREDICTION

*Slides in this section are by Lieven Eeckhout, Ghent University.  
Reused with permission.*

# Branch target prediction

- Branch target buffer (BTB) or Branch target address cache (BTAC)
- Small (set-associative) cache
- Idea:
  - Input = branch address
  - Output = branch target address
  - Accessed in the same cycle as the I-cache access
  - Branch target address is the fetch address in the next cycle
- Done for both conditional as well as non-conditional branches


# Branch Target Buffer




# Return Address Stack

- Function calls are easy to predict
- Function returns are much more complicated to predict
- Solution: Return Address Stack (RAS)
  - by Kaeli and Emma in 1991
  - In Pentium 4: 16 entries
  - Circular buffer
    - Push function return address on RAS; pop the address upon a return
    - If depth of function call stack > RAS size: incorrect predictions

# Return Address Stack



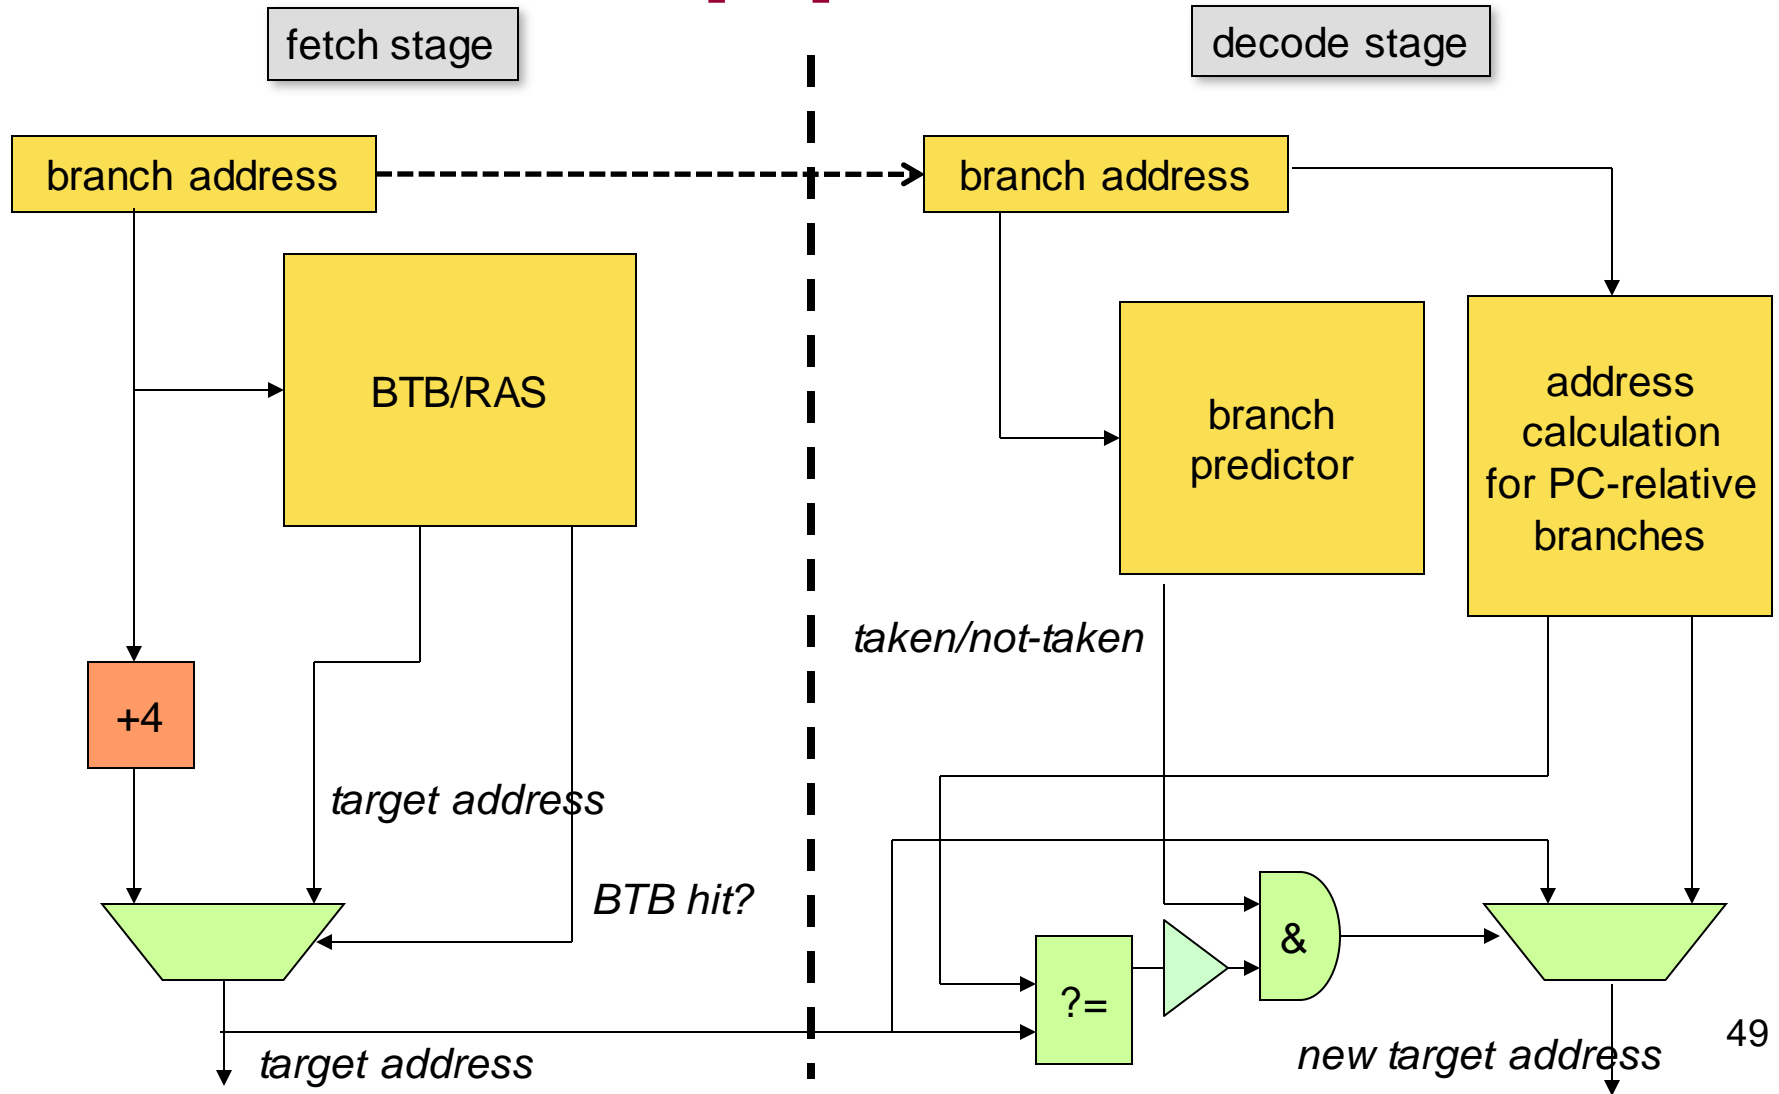
In case of a function call



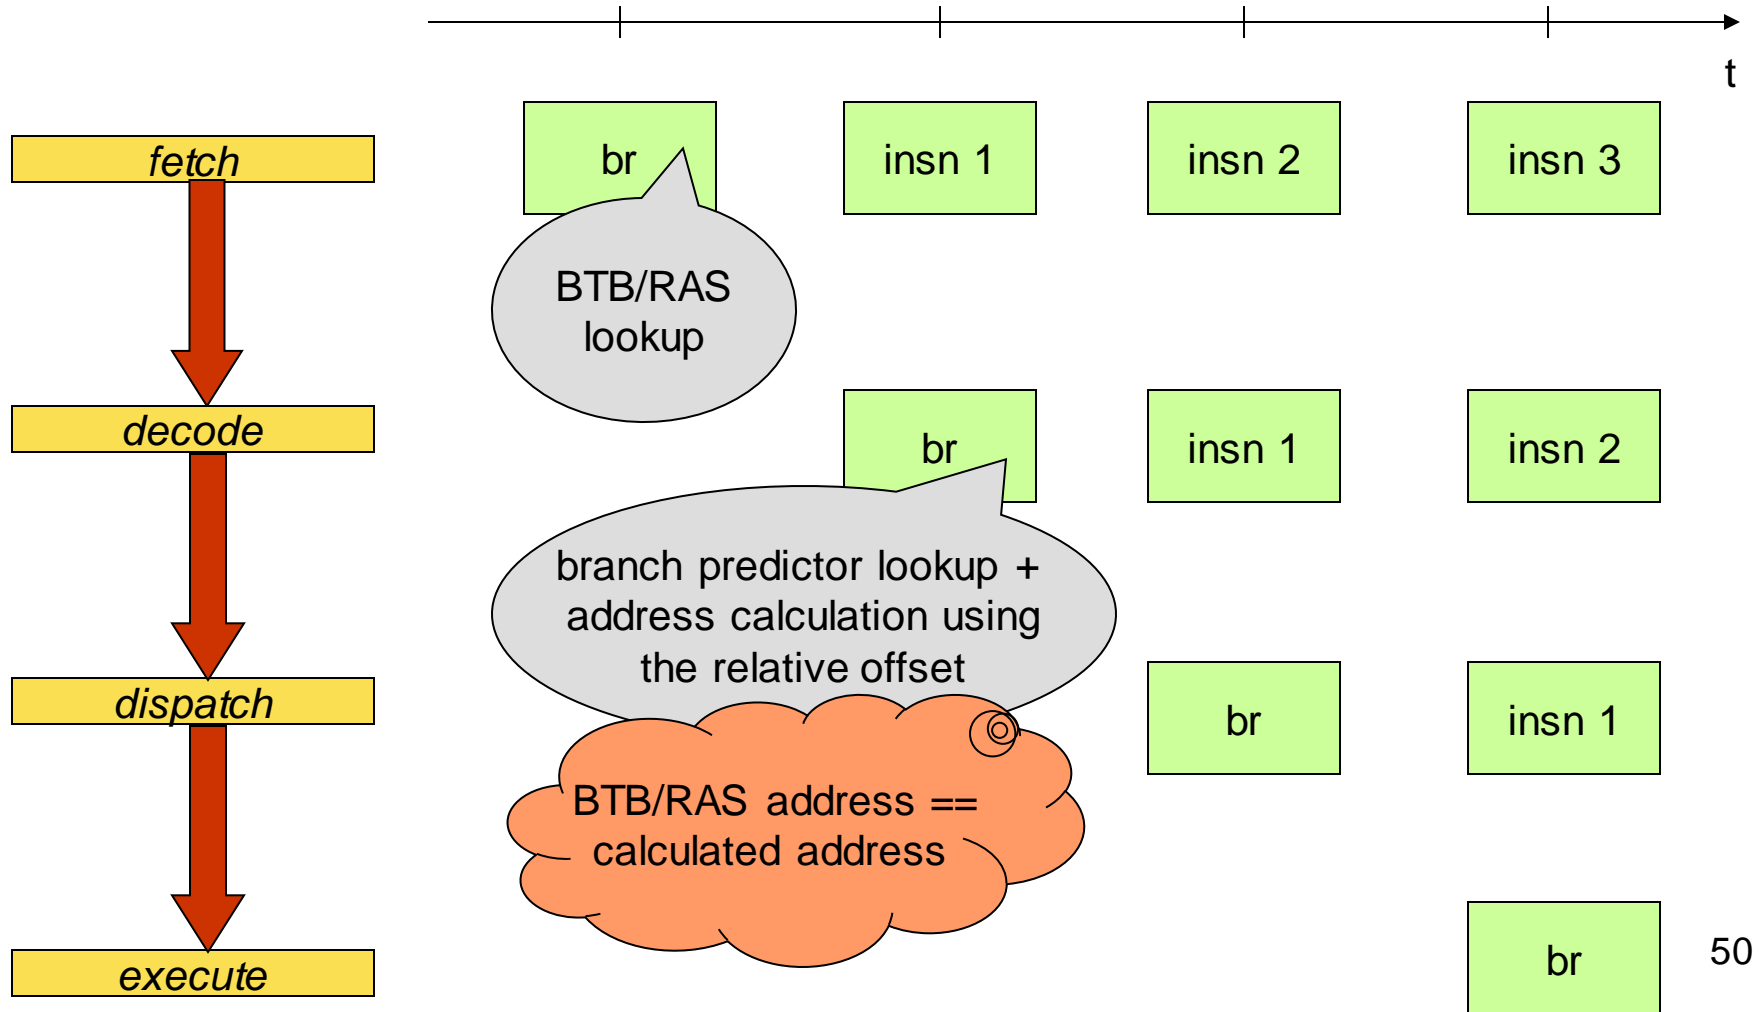
In case of a return



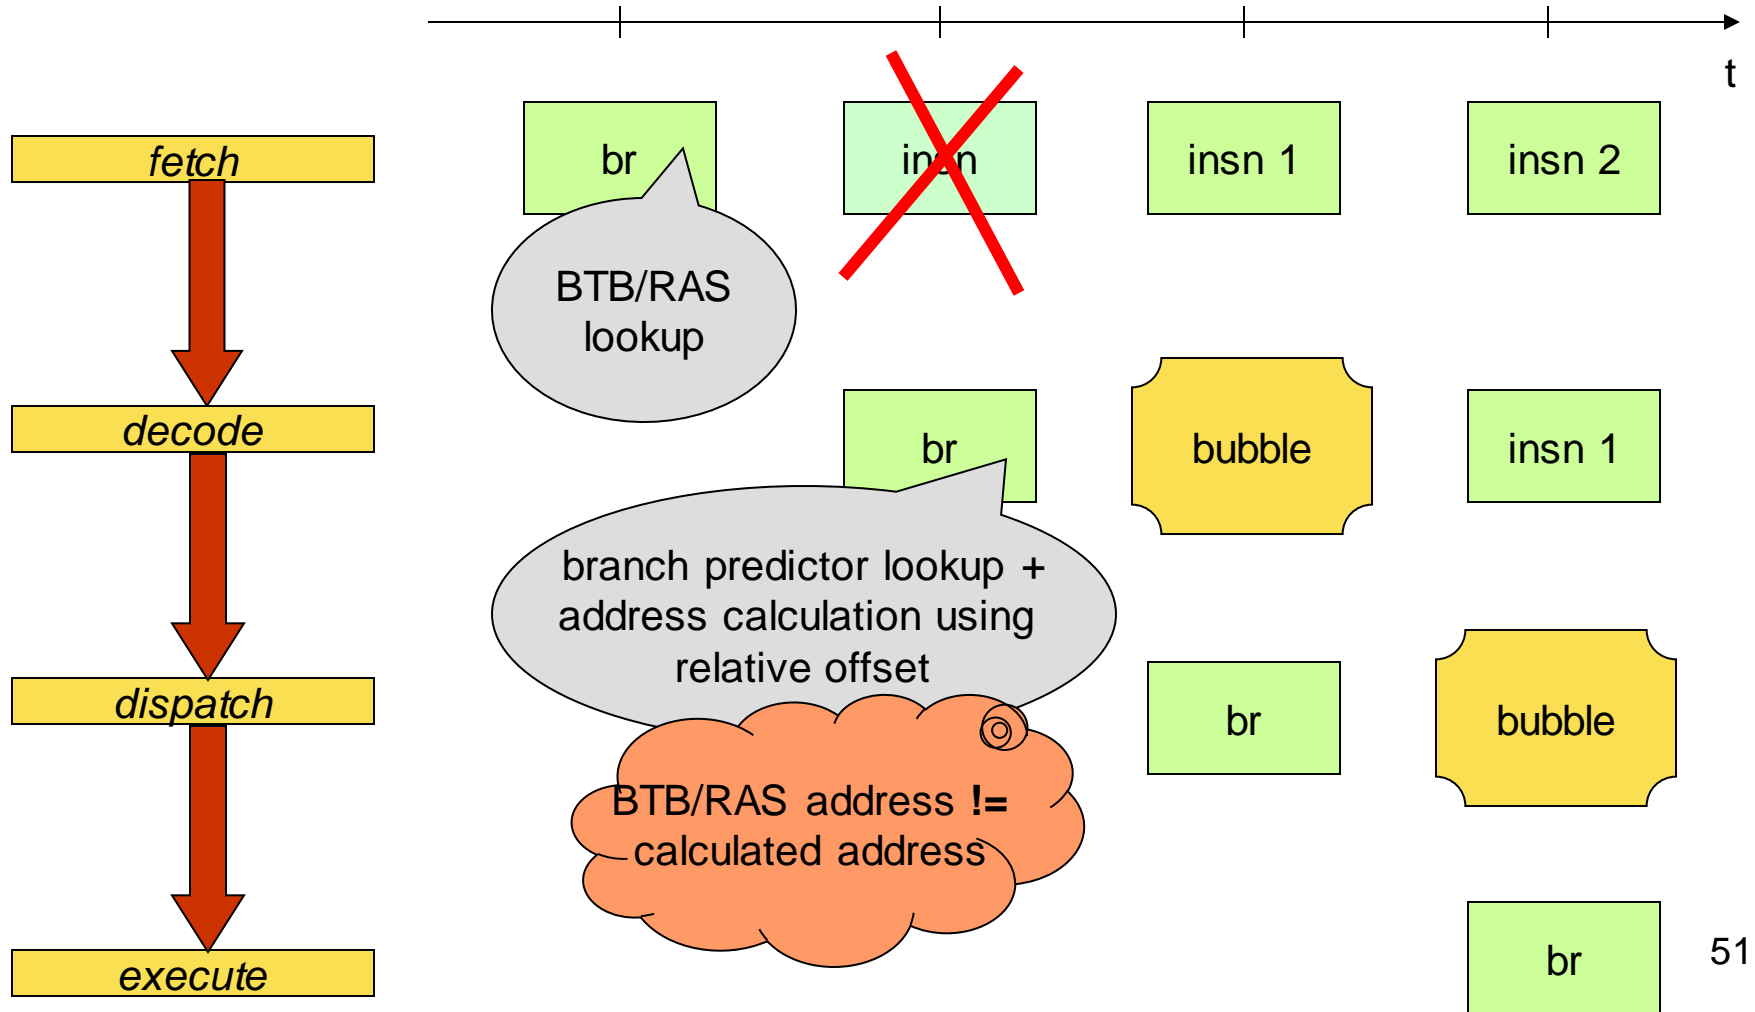
# Branch prediction in the pipeline



# Branch prediction predicts taken and target address is correct



# Branch predictor predicts taken but target address is incorrect




# BRANCH MISPREDICTION RECOVERY

*Slides in this section are by Lieven Eeckhout, Ghent University.  
Reused with permission.*

# Speculative execution

- Predict branch target address and start fetching and executing instructions along the predicted path
  - no completion of speculative insns!
- There might be multiple branches in flight along the predicted path
  - tags are added


# Speculative execution



# Pipeline squashing


- Branch outcome is known at execution stage
- If correctly predicted
  - Corresponding tags are deallocated
  - Instructions become non-speculative
- If mispredicted
  - Wrong-path instructions are nullified
  - Start fetching instructions along correct path

# Speculative execution





# Squash wrong-path insns and fetch insns along correct path



# SUMMARY

# Summary

- Branch prediction is critical to achieving high performance
  - Must be able to fill the pipeline with (mostly) useful instructions
  - The deeper the pipeline, the more important branch prediction becomes
  - Luckily, branches are quite predictable (90-99% prediction accuracy is common)
- Branch prediction schemes
  - Statically predict branch direction: Rule-based, program-based, and profile-based
  - Dynamically predict branch direction: Bimodal predictors, two-level predictors, and tournament predictors
  - Predict the branch target: Branch Target Buffer (BTB) and Return Address Stack (RAS)